



OpenTofu

iLGlabs

Redefining how you think, learn, and work.

iLGlabs

By : JP

info@gnugroup.org

www.gnugroup.org

Introduction to OpenTofu {#introduction}

OpenTofu is an open-source Infrastructure as Code (IaC) tool that enables you to define, provision, and manage infrastructure using declarative configuration files. It was forked from Terraform in 2023 and is maintained by the Linux Foundation.

Why OpenTofu?

- **Open Source:** Truly open-source with community governance
- **Declarative:** Describe your desired infrastructure state
- **Cloud Agnostic:** Works with AWS, Azure, GCP, and 100+ providers
- **Version Control:** Infrastructure as code in Git
- **State Management:** Tracks actual vs desired state

Core Concepts and Terminology {#core-concepts}

1. Provider

A provider is a plugin that interacts with a specific cloud platform or service API.

```
hcl

provider "aws" {
  region = "us-east-1"
}
```

2. Resource

A resource is a component of your infrastructure (VM, database, network, etc.).

```
hcl

resource "aws_instance" "example" {
    ami          = "ami-0c55b159cbfafe1f0"
    instance_type = "t2.micro"
}
```

3. Data Source

A data source allows you to fetch information from providers without creating resources.

```
hcl

data "aws_ami" "ubuntu" {
    most_recent = true
    owners       = ["099720109477"]
}
```

4. Variable

Variables allow parameterization of your configuration.

```
hcl
```

```
variable "instance_type" {
    description = "EC2 instance type"
    type        = string
    default     = "t2.micro"
}
```

5. Output

Outputs display information after infrastructure is created.

```
hcl

output "instance_ip" {
    value = aws_instance.example.public_ip
}
```

6. State

State is a file that maps your configuration to real-world resources. OpenTofu uses this to track what it manages.

7. Module

A module is a container for multiple resources used together. It promotes reusability.

8. Backend

A backend determines where state files are stored (local, S3, Azure Blob, etc.).

Providers, Resources & Data Sources

Learning Objectives

By completing this lab, you will:

- Understand provider configuration and version constraints
- Configure multiple providers in a single project
- Create resources with implicit and explicit dependencies
- Use data sources to query existing infrastructure
- Link resources dynamically using data source outputs
- Apply explicit dependency management with `depends_on`

Part 1: Conceptual Overview

1.1 Provider Sources and Versions

What is a Provider?

A provider is a plugin that enables OpenTofu to interact with APIs of cloud platforms, SaaS services, or other infrastructure.

Providers are responsible for:

Understanding API interactions

Exposing resources and data sources

Managing authentication

Provider Configuration Elements:

```
hcl

terraform {
  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 5.0"          # Registry namespace/name
                                    # Version constraint
    }
  }

  provider "aws" {
    region = var.aws_region      # Provider-specific configuration
  }
}
```

Version Constraint Operators:

- `= 1.0.0` - Exact version
- `>= 1.0.0` - Greater than or equal to
- `<= 1.0.0` - Less than or equal to
- `~> 1.0.0` - Allows only rightmost version increment (1.0.x)
- `~> 1.0` - Allows minor version updates (1.x)

1.2 Multiple Provider Usage

Why Use Multiple Providers?

Deploy across multiple cloud platforms

Use utility providers (random, time, null)

Deploy to multiple regions of the same cloud

Integrate with various SaaS services

Provider Aliases:

hcl

```
provider "aws" {
  region = "us-east-1"
  alias   = "primary"
}

provider "aws" {
  region = "us-west-2"
  alias   = "secondary"
}

resource "aws_instance" "primary_server" {
  provider = aws.primary
  # ... configuration
}
```

1.3 Resources and Dependencies

Resources are the infrastructure objects you want to create/manage:

- EC2 instances
- S3 buckets
- Security groups
- Load balancers

Implicit Dependencies (Automatic):

OpenTofu automatically detects dependencies when you reference one resource in another:

```
resource "aws_security_group" "app_sg" {
    # ... configuration
}

resource "aws_instance" "app_server" {
    security_groups = [aws_security_group.app_sg.id] # Implicit dependency
}
```



Explicit Dependencies (Manual): Use `depends_on` when OpenTofu can't detect the relationship:

```
resource "aws_iam_role" "app_role" {
    # ... configuration
}

resource "aws_instance" "app_server" {
    depends_on = [aws_iam_role.app_role] # Explicit dependency
}
```

1.4 Data Sources

What are Data Sources?

Data sources allow OpenTofu to fetch information about existing infrastructure without managing it.

They're read-only queries.

Common Use Cases:

Look up the latest AMI

Get current AWS account information

Query existing VPC or subnet details

Retrieve DNS information

```
data "aws_ami" "amazon_linux" {
    most_recent = true

    filter {
        name    = "name"
        values = ["al2023-ami--x86_64"]
    }

    owners = ["amazon"]
}

# Reference in resource
resource "aws_instance" "server" {
    ami = data.aws_ami.amazon_linux.id # Dynamic AMI reference
}
...
```

Lab Scenario

You'll create a multi-provider infrastructure that:

1. Uses AWS provider for cloud resources
2. Uses random provider for generating unique identifiers
3. Queries existing AWS infrastructure with data sources
4. Creates an EC2 instance with proper dependencies
5. Demonstrates both implicit and explicit dependencies

Prerequisites Checklist

- [] OpenTofu installed (version **1.6+**)
- [] AWS CLI configured with credentials
- [] AWS account with EC2 permissions
- [] Text editor (VS Code recommended)
- [] Git installed (optional, for version control)

```
### Project Structure
```
multi-provider-lab/
├── versions.tf # Provider requirements and versions
├── variables.tf # Input variables
├── main.tf # Primary resource definitions
├── data.tf # Data source queries
├── outputs.tf # Output values
└── README.md # Lab documentation
```

lesson8\_multi\_provider\_lab

# **OpenTofu : Variables, Functions & Locals**

# Module Overview

This training module covers advanced OpenTofu variable handling, built-in functions, and dynamic block generation.

We'll learn how to create flexible, reusable infrastructure code that adapts to different requirements without code duplication.

## Learning Objectives:

Master complex variable types for sophisticated data structures

Use local values for computed and derived values

Leverage built-in functions for data transformation

Implement dynamic blocks for conditional resource generation

Build a parameterized module with dynamic resource creation

# 1. Complex Variable Types

## 1.1 Understanding Type Constraints

OpenTofu supports several complex variable types beyond simple strings and numbers:

### Basic Types:

- `string` - Text values
- `number` - Numeric values (integers and floats)
- `bool` - True/false values

### Complex Types:

- `list(type)` - Ordered collection of values
- `map(type)` - Key-value pairs with string keys
- `set(type)` - Unordered collection of unique values
- `object({})` - Complex structure with named attributes
- `tuple([])` - Ordered collection with specific types

## 1.2 Maps

Maps store key-value pairs where keys are strings. Useful for configurations that vary by environment, region, or other categories.

```
hcl

variable "instance_types" {
 description = "EC2 instance types by environment"
 type = map(string)
 default = {
 dev = "t3.micro"
 staging = "t3.small"
 prod = "t3.large"
 }
}

Usage
resource "aws_instance" "app" {
 instance_type = var.instance_types[var.environment]
 # ...
}
```

## Map of Maps (Nested):

---

```
hcl

variable "regional_settings" {
 description = "Settings by region"
 type = map(map(string))
 default = {
 us-east-1 = {
 ami = "ami-0c55b159cbfafe1f0"
 instance_type = "t3.small"
 }
 us-west-2 = {
 ami = "ami-0d1cd67c26f5fca19"
 instance_type = "t3.medium"
 }
 }
}

Usage
resource "aws_instance" "regional" {
 ami = var.regional_settings[var.region]["ami"]
 instance_type = var.regional_settings[var.region]["instance_type"]
}
```

## 1.3 Objects

Objects define complex structures with specific attribute types. Perfect for grouping related configuration.

hcl

```
variable "database_config" {
 description = "Database configuration"
 type = object({
 engine = string
 engine_version = string
 instance_class = string
 allocated_storage = number
 multi_az = bool
 backup_retention = number
 })
 default = {
 engine = "mysql"
 engine_version = "8.0"
 instance_class = "db.t3.micro"
 allocated_storage = 20
 multi_az = false
 backup_retention = 7
 }
}

Usage
resource "aws_db_instance" "main" {
 engine = var.database_config.engine
 engine_version = var.database_config.engine_version
 instance_class = var.database_config.instance_class
 allocated_storage = var.database_config.allocated_storage
 multi_az = var.database_config.multi_az
 backup_retention_period = var.database_config.backup_retention
}
```

## List of Objects (Very Common Pattern):

hcl

```
variable "subnets" {
 description = "List of subnet configurations"
 type = list(object({
 name = string
 cidr_block = string
 availability_zone = string
 public = bool
 tags = map(string)
 }))
 default = [
 {
 name = "public-subnet-1"
 cidr_block = "10.0.1.0/24"
 availability_zone = "us-east-1a"
 public = true
 tags = { Tier = "Public" }
 },
 {
 name = "private-subnet-1"
 cidr_block = "10.0.10.0/24"
 availability_zone = "us-east-1a"
 public = false
 tags = { Tier = "Private" }
 }
]
}
```

```
Usage with for_each
resource "aws_subnet" "main" {
 for_each = { for subnet in var.subnets : subnet.name => subnet }

 vpc_id = aws_vpc.main.id
 cidr_block = each.value.cidr_block
 availability_zone = each.value.availability_zone
 map_public_ip_on_launch = each.value.public

 tags = merge(
 each.value.tags,
 { Name = each.value.name }
)
}
```

## 1.4 Tuples

Tuples are ordered collections with specific types for each element. Less common but useful for fixed-structure data.

---

```
hcl

variable "server_config" {
 description = "Server configuration [name, instance_type, count]"
 type = tuple([string, string, number])
 default = ["web-server", "t3.micro", 3]
}

Usage
locals {
 server_name = var.server_config[0]
 instance_type = var.server_config[1]
 server_count = var.server_config[2]
}
```

## 1.5 Sets

Sets are unordered collections of unique values. Useful when order doesn't matter and duplicates should be eliminated.

```
hcl

variable "allowed_cidr_blocks" {
 description = "Allowed CIDR blocks for security group"
 type = set(string)
 default = ["10.0.0.0/8", "172.16.0.0/12", "192.168.0.0/16"]
}

Usage
resource "aws_security_group_rule" "allow_internal" {
 for_each = var.allowed_cidr_blocks

 type = "ingress"
 from_port = 443
 to_port = 443
 protocol = "tcp"
 cidr_blocks = [each.value]
 security_group_id = aws_security_group.main.id
}
```

## 2. Local Values and Computed Variables

### 2.1 What are Local Values?

Local values are named expressions that you can reference throughout your configuration.

They help:

- Avoid repetition
- Compute derived values
- Make complex expressions more readable
- Combine multiple variable values

#### Syntax:

```
hcl
locals {
 name = expression
}

Reference with local.name
```

## 2.2 Common Local Value Patterns

### Environment-specific naming:

hcl

```
variable "project_name" {
 type = string
 default = "myapp"
}

variable "environment" {
 type = string
 default = "dev"
}

locals {
 # Compute a common prefix for all resources
 name_prefix = "${var.project_name}-${var.environment}"

 # Common tags applied to all resources
 common_tags = {
 Project = var.project_name
 Environment = var.environment
 ManagedBy = "OpenTofu"
 CreatedAt = timestamp()
 }
}
```

```
Usage
resource "aws_vpc" "main" {
 cidr_block = "10.0.0.0/16"

 tags = merge(
 local.common_tags,
 {
 Name = "${local.name_prefix}-vpc"
 }
)
}
```

## Computing CIDR blocks:

---

```
hcl

variable "vpc_cidr" {
 type = string
 default = "10.0.0.0/16"
}

locals {
 # Split VPC CIDR into subnet ranges
 subnet_count = 3
 subnet_newbits = 8 # Creates /24 subnets from /16

 public_subnet_cidrs = [
 for i in range(local.subnet_count) :
 cidrsubnet(var.vpc_cidr, local.subnet_newbits, i)
]

 private_subnet_cidrs = [
 for i in range(local.subnet_count) :
 cidrsubnet(var.vpc_cidr, local.subnet_newbits, i + 10)
]
}

Result:
public_subnet_cidrs = ["10.0.0.0/24", "10.0.1.0/24", "10.0.2.0/24"]
private_subnet_cidrs = ["10.0.10.0/24", "10.0.11.0/24", "10.0.12.0/24"]
```

## Flattening nested structures:

```
hcl

variable "applications" {
 type = map(object({
 subnets = list(string)
 ports = list(number)
 }))
 default = {
 web = {
 subnets = ["subnet-1", "subnet-2"]
 ports = [80, 443]
 }
 api = {
 subnets = ["subnet-3"]
 ports = [8080, 8443]
 }
 }
}

locals {
 # Flatten to create one security group rule per app per port
 security_rules = flatten([
 for app_name, app in var.applications : [
 for port in app.ports : {
 app_name = app_name
 port = port
 }
]
])
}

Result:
[
{app_name = "web", port = 80},
{app_name = "web", port = 443},
{app_name = "api", port = 8080},
{app_name = "api", port = 8443}
]
```

## 2.3 Local Values vs Variables

### **Use Variables when:**

- Values should be configurable by users
- Different values needed per environment
- Values come from external sources

### **Use Locals when:**

- Computing derived values
- Complex expressions used multiple times
- Internal logic not meant for external configuration

### **3. Built-in Functions**

OpenTofu provides 100+ built-in functions for data transformation.

### 3.1 Lookup Functions

#### **lookup(map, key, default)**

---

hcl

```
variable "instance_types" {
 type = map(string)
 default = {
 dev = "t3.micro"
 prod = "t3.large"
 }
}

Safely lookup with fallback
locals {
 instance_type = lookup(var.instance_types, var.environment, "t3.small")
}

If var.environment = "staging" (not in map), returns "t3.small"
```

## element(list, index)

---

hcl

```
variable "availability_zones" {
 type = list(string)
 default = ["us-east-1a", "us-east-1b", "us-east-1c"]
}

Cycle through AZs for multiple resources
resource "aws_subnet" "example" {
 count = 6
 availability_zone = element(var.availability_zones, count.index)
 # Creates 2 subnets per AZ: 1a, 1b, 1c, 1a, 1b, 1c
}
```

## index(list, value)

```
hcl
```

```
locals {
 environments = ["dev", "staging", "prod"]
 env_index = index(local.environments, var.environment)
 # Returns 0 for dev, 1 for staging, 2 for prod
}
```

## 3.2 File Functions

### **file(path)**

hcl

```
Read file contents as string
resource "aws_instance" "web" {
 user_data = file("${path.module}/scripts/init.sh")
}
```

## **templatefile(path, vars)**

hcl

```
Read and interpolate template
resource "aws_instance" "web" {
 user_data = templatefile("${path.module}/templates/init.sh.tpl", {
 db_host = aws_db_instance.main.endpoint
 environment = var.environment
 app_version = var.app_version
 })
}

init.sh.tpl:
#!/bin/bash
export DB_HOST="${db_host}"
export ENVIRONMENT="${environment}"
export APP_VERSION="${app_version}"
```

## filebase64(path)

---

hcl

```
Read file as base64 (useful for binary files)
resource "aws_s3_object" "image" {
 bucket = aws_s3_bucket.assets.id
 key = "logo.png"
 content_base64 = filebase64("${path.module}/assets/logo.png")
}
```

### 3.3 String Functions

#### **format(spec, values...)**

```
hcl
```

```
locals {
 instance_name = format("%s-%s-instance-%02d",
 var.project_name,
 var.environment,
 count.index + 1
)
 # Result: "myapp-prod-instance-01"
}
```

## join(separator, list)

---

hcl

```
variable "name_parts" {
 type = list(string)
 default = ["myapp", "web", "server"]
}

locals {
 full_name = join("-", var.name_parts)
 # Result: "myapp-web-server"
}
```

## split(separator, string)

```
hcl
```

```
locals {
 db_endpoint = "mydb.abc123.us-east-1.rds.amazonaws.com:3306"
 db_parts = split(":", local.db_endpoint)
 db_host = local.db_parts[0]
 db_port = local.db_parts[1]
}
```

## **replace(string, search, replace)**

```
hcl
```

```
locals {
 sanitized_name = replace(var.project_name, "_", "-")
 # Replaces underscores with hyphens
}
```

### 3.4 Numeric Functions

#### **min(numbers...) / max(numbers...)**

---

hcl

```
variable "instance_counts" {
 type = map(number)
 default = {
 web = 3
 api = 5
 worker = 2
 }
}

locals {
 min_instances = min(var.instance_counts["web"],
 var.instance_counts["api"],
 var.instance_counts["worker"])
 max_instances = max(var.instance_counts...) # Splat syntax
}
```

## **ceil(number) / floor(number)**

---

hcl

```
locals {
 storage_gb_needed = 47.8
 storage_rounded = ceil(local.storage_gb_needed) # 48
}
```

### 3.5 Collection Functions

#### `merge(maps...)`

```
hcl

locals {
 default_tags = {
 ManagedBy = "OpenTofu"
 Project = var.project_name
 }

 environment_tags = {
 Environment = var.environment
 }

 all_tags = merge(
 local.default_tags,
 local.environment_tags,
 var.custom_tags
)
}
```

## **concat(lists...)**

hcl

```
locals {
 public_subnets = ["10.0.1.0/24", "10.0.2.0/24"]
 private_subnets = ["10.0.10.0/24", "10.0.11.0/24"]
 all_subnets = concat(local.public_subnets, local.private_subnets)
}
```

## **flatten(list\_of\_lists)**

---

hcl

```
locals {
 nested_cidrs = [
 ["10.0.1.0/24", "10.0.2.0/24"],
 ["10.0.10.0/24", "10.0.11.0/24"]
]
 flat_cidrs = flatten(local.nested_cidrs)
 # ["10.0.1.0/24", "10.0.2.0/24", "10.0.10.0/24", "10.0.11.0/24"]
}
```

## keys(map) / values(map)

```
hcl
```

```
variable "regions" {
 type = map(string)
 default = {
 us-east-1 = "N. Virginia"
 us-west-2 = "Oregon"
 }
}

locals {
 region_codes = keys(var.regions) # ["us-east-1", "us-west-2"]
 region_names = values(var.regions) # ["N. Virginia", "Oregon"]
}
```

### 3.6 Conditional Functions

#### **coalesce(values...)**

hcl

```
Returns first non-null, non-empty value
locals {
 instance_type = coalesce(
 var.instance_type_override,
 var.default_instance_type,
 "t3.micro"
)
 # Uses override if set, otherwise default, otherwise t3.micro
}
```

## try(expressions...)

```
hcl
```

```
Returns first expression that doesn't error
locals {
 ami_id = try(
 var.custom_ami_id,
 data.aws_ami.amazon_linux.id,
 "ami-default"
)
}
```

## **can(expression)**

---

hcl

```
Returns true if expression evaluates without error
locals {
 has_custom_ami = can(regex("^\ ami-", var.ami_id))
}
```

## 3.7 Network Functions

### **cidrsubnet(prefix, newbits, netnum)**

---

hcl

```
Split VPC CIDR into smaller subnets
locals {
 vpc_cidr = "10.0.0.0/16"

 # Create /24 subnets
 subnet_1 = cidrsubnet(local.vpc_cidr, 8, 0) # 10.0.0.0/24
 subnet_2 = cidrsubnet(local.vpc_cidr, 8, 1) # 10.0.1.0/24
 subnet_3 = cidrsubnet(local.vpc_cidr, 8, 10) # 10.0.10.0/24
}
```

## **cidrhost(prefix, hostnum)**

hcl

```
Get specific IP from CIDR block
locals {
 subnet = "10.0.1.0/24"
 first_ip = cidrhost(local.subnet, 1) # 10.0.1.1
 tenth_ip = cidrhost(local.subnet, 10) # 10.0.1.10
}
```

## 4. Dynamic and Conditional Blocks

### 4.1 Dynamic Blocks

Dynamic blocks allow you to generate nested blocks within resources based on variable content. Essential for creating flexible, reusable modules.

#### Basic Syntax:

---

```
hcl

dynamic "block_name" {
 for_each = var.list_or_map
 content {
 # Block attributes using each.value or each.key
 }
}
```

## Example: Dynamic Ingress Rules

```
hcl

variable "ingress_rules" {
 description = "List of ingress rules"
 type = list(object({
 from_port = number
 to_port = number
 protocol = string
 cidr_blocks = list(string)
 description = string
 }))
 default = [
 {
 from_port = 80
 to_port = 80
 protocol = "tcp"
 cidr_blocks = ["0.0.0.0/0"]
 description = "HTTP from anywhere"
 },
 {
 from_port = 443
 to_port = 443
 protocol = "tcp"
 cidr_blocks = ["0.0.0.0/0"]
 description = "HTTPS from anywhere"
 }
]
}

resource "aws_security_group" "web" {
 name = "web-sg"
 vpc_id = aws_vpc.main.id

 dynamic "ingress" {
 for_each = var.ingress_rules
 content {
 from_port = ingress.value.from_port
 to_port = ingress.value.to_port
 protocol = ingress.value.protocol
 cidr_blocks = ingress.value.cidr_blocks
 description = ingress.value.description
 }
 }
}
```

## Example: Dynamic Tags

```
hcl

variable "enable_detailed_monitoring" {
 type = bool
 default = false
}

resource "aws_instance" "web" {
 ami = data.aws_ami.amazon_linux.id
 instance_type = "t3.micro"

 dynamic "tag_specifications" {
 for_each = var.enable_detailed_monitoring ? [1] : []
 content {
 resource_type = "instance"
 tags = {
 DetailedMonitoring = "enabled"
 }
 }
 }
}
```

## 4.2 Conditional Blocks with for\_each

Use empty maps or sets to conditionally create blocks:

```
hcl

variable "enable_backup" {
 type = bool
 default = true
}

resource "aws_db_instance" "main" {
 # ... other config ...

 dynamic "backup_configuration" {
 for_each = var.enable_backup ? [1] : []
 content {
 backup_retention_period = 7
 backup_window = "03:00-04:00"
 maintenance_window = "mon:04:00-mon:05:00"
 }
 }
}
```

## 4.3 Nested Dynamic Blocks

You can nest dynamic blocks for complex structures:

```
hcl

variable "load_balancer_listeners" {
 type = list(object({
 port = number
 protocol = string
 rules = list(object({
 priority = number
 host = string
 path = string
 }))
 }))
}

resource "aws_lb_listener" "main" {
 for_each = { for listener in var.load_balancer_listeners : listener.port => listener }

 load_balancer_arn = aws_lb.main.arn
 port = each.value.port
 protocol = each.value.protocol

 dynamic "default_action" {
 for_each = each.value.rules
 content {
 type = "forward"
 }
 }

 dynamic "forward" {
 for_each = [default_action.value]
 content {
 target_group {
 arn = aws_lb_target_group.main.arn
 weight = 100
 }
 }
 }
}
```

## 4.4 Conditional Resource Creation

Using count:

```
hcl

variable "create_nat_gateway" {
 type = bool
 default = true
}

resource "aws_nat_gateway" "main" {
 count = var.create_nat_gateway ? 1 : 0

 allocation_id = aws_eip.nat[0].id
 subnet_id = aws_subnet.public[0].id
}

Reference with [0] when used
resource "aws_route" "private_nat" {
 count = var.create_nat_gateway ? 1 : 0

 route_table_id = aws_route_table.private.id
 destination_cidr_block = "0.0.0.0/0"
 nat_gateway_id = aws_nat_gateway.main[0].id
}
```

## Using `for_each`:

---

hcl

```
variable "create_bastion" {
 type = bool
 default = false
}

resource "aws_instance" "bastion" {
 for_each = var.create_bastion ? { bastion = true } : {}

 ami = data.aws_ami.amazon_linux.id
 instance_type = "t3.micro"
 subnet_id = aws_subnet.public[0].id
}
```

## Key Takeaways

1. **Complex Types** enable sophisticated data structures matching real-world infrastructure needs
2. **Local Values** reduce repetition and make complex logic readable
3. **Built-in Functions** provide powerful data transformation without external dependencies
4. **Dynamic Blocks** create flexible modules that adapt to variable inputs
5. **Conditional Logic** allows single codebase to handle multiple scenarios

## **LAB: Parameterized VPC Module with Dynamic Subnets and Security Groups**

# Lab Overview

Build a reusable VPC module that dynamically creates subnets across availability zones and generates security groups with flexible rule sets. This lab demonstrates all advanced variable concepts in a real-world scenario.

## What You'll Build:

Parameterized VPC module accepting complex variable types

Dynamic subnet creation across multiple AZs

Dynamic security group generation with computed rules

Local values for CIDR calculations

Built-in functions for data transformation

## Architecture

---

### VPC Module

- └ Dynamic Public Subnets (across multiple AZs)
  - └ Dynamic Private Subnets (across multiple AZs)
  - └ Dynamic Database Subnets (across multiple AZs)
  - └ Internet Gateway (conditional)
  - └ NAT Gateways (conditional, per AZ)
  - └ Route Tables (dynamic based on subnet tiers)
  - └ Security Groups (dynamic rules based on application requirements)
-

## Project Structure

Create this directory structure:

```
lab-vpc-module/
 └── modules/
 └── vpc/
 ├── main.tf
 ├── variables.tf
 ├── outputs.tf
 ├── locals.tf
 ├── subnets.tf
 ├── security-groups.tf
 └── routing.tf
 └── environments/
 └── dev/
 ├── main.tf
 ├── variables.tf
 ├── terraform.tfvars
 └── versions.tf
 └── prod/
 ├── main.tf
 ├── variables.tf
 ├── terraform.tfvars
 └── versions.tf
 └── README.md
```

LAB - lesson11\_Adv\_variables\_functions\_locals\_vpc\_modules

# **OpenTofu Modules and Registries**

## **Learning Objectives**

By the end of this training, you will be able to:

- Understand the difference between root and child modules
- Create reusable, production-ready modules
- Use local and remote module sources effectively
- Leverage community modules from registries
- Publish and version your own modules
- Implement proper module versioning strategies
- Apply best practices for module design and composition

# 1. Understanding OpenTofu Modules

## What Are Modules?

A **module** is a container for multiple resources that are used together.

Every OpenTofu configuration has at least one module, called the **root module**, which consists of the resources defined in the `.tf` files in the main working directory.

## **Root vs Child Modules**

### **Root Module:**

The main working directory where you run `tofu` commands  
Contains your primary configuration files  
Calls child modules to compose infrastructure  
Manages state directly

### **Child Modules:**

Reusable components called by the root module  
Self-contained with their own resources, variables, and outputs  
Can be called multiple times with different configurations  
Promote DRY (Don't Repeat Yourself) principle

## Why Use Modules?

1. **Reusability:** Write once, use many times
2. **Organization:** Logical grouping of related resources
3. **Encapsulation:** Hide complexity behind simple interfaces
4. **Standardization:** Enforce organizational best practices
5. **Collaboration:** Share infrastructure patterns across teams
6. **Testing:** Isolate and test components independently

## 2. Creating Reusable Modules

### Module Structure Best Practices

A well-structured module typically contains:

```
module-name/
├── README.md # Documentation
├── main.tf # Primary resource definitions
├── variables.tf # Input variables
├── outputs.tf # Output values
├── versions.tf # Provider requirements
└── examples/
 └── basic/
 ├── main.tf
 └── variables.tf
└── tests/ # Validation tests (optional)
```

## **Key Principles for Module Design**

## **1. Single Responsibility**

- Each module should do one thing well
- Focus on a specific infrastructure pattern
- Avoid overly complex, multi-purpose modules

## **2. Configurable But Opinionated**

- Provide sensible defaults
- Allow customization through variables
- Guide users toward best practices

### **3. Clear Interface**

- Well-documented input variables
- Meaningful output values
- Explicit requirements and constraints

### **4. Composability**

- Modules should work well together
- Avoid tight coupling between modules
- Use outputs/inputs for module communication

## Example: Simple Module Structure

### Root Module (calls the module):

```
hcl

module "web_server" {
 source = "./modules/ec2-instance"

 instance_name = "web-server-1"
 instance_type = "t3.micro"
 subnet_id = aws_subnet.public.id
}
```

### Child Module (../modules/ec2-instance/main.tf):

```
hcl

resource "aws_instance" "this" {
 ami = var.ami_id
 instance_type = var.instance_type
 subnet_id = var.subnet_id

 tags = {
 Name = var.instance_name
 }
}
```

### **3. Module Sources: Local vs Remote**

### 3. Module Sources: Local vs Remote

#### Local Modules

##### Path Syntax:

hcl

```
module "local_example" {
 source = "./modules/vpc" # Relative path
 source = "/absolute/path/vpc" # Absolute path (not recommended)
}
```

##### Use Cases:

- Development and testing
- Organization-specific modules not for public use
- Tightly coupled components
- When you need to modify modules frequently

##### Advantages:

- Fast iteration during development
- No versioning complexity
- Full control over code

##### Disadvantages:

- No version control
- Must be distributed with root module
- Harder to share across projects

## 1. Git Repositories

### Remote Modules

OpenTofu supports multiple remote sources:

```
HTTPS
module "vpc" {
 source = "git::https://github.com/example/terraform-aws-vpc.git"
}

SSH
module "vpc" {
 source = "git::ssh://git@github.com/example/terraform-aws-vpc.git"
}

Specific branch
module "vpc" {
 source = "git::https://github.com/example/terraform-aws-vpc.git?ref=main"
}

Specific tag (versioning)
module "vpc" {
 source = "git::https://github.com/example/terraform-aws-vpc.git?ref=v1.2.0"
}

Specific commit
module "vpc" {
 source = "git::https://github.com/example/terraform-aws-vpc.git?ref=abc123"
}

Subdirectory
module "vpc" {
 source = "git::https://github.com/example/modules.git//vpc"
}
```

## 2. HTTP/HTTPS URLs

### 2. HTTP/HTTPS URLs

---

hcl

```
module "vpc" {
 source = "https://example.com/modules/vpc.zip"
}
```

### 3. Registry Sources (Terraform Registry)

```
hcl

Terraform Registry format
module "vpc" {
 source = "terraform-aws-modules/vpc/aws"
 version = "5.1.0"
}
```

### 4. Local Archives

```
hcl

module "vpc" {
 source = "./modules/vpc-1.0.0.zip"
}
```

## Module Versioning Best Practices

### Semantic Versioning (SemVer):

- Format: **MAJOR.MINOR.PATCH** (e.g., **1.4.2**)
- **MAJOR**: Breaking changes
- **MINOR**: New features, backward compatible
- **PATCH**: Bug fixes, backward compatible

### Version Constraints:

```
hcl

module "vpc" {
 source = "git::https://github.com/example/vpc.git?ref=v1.2.0"
 # Or for registry modules:
 version = "~> 1.2.0" # >= 1.2.0, < 1.3.0
 version = ">= 1.2.0" # Any version >= 1.2.0
 version = "1.2.0" # Exact version
}
```

## 4. Using Community Registry Modules

### OpenTofu Registry (Library.tf)

**Library.tf** is the community-driven module registry for OpenTofu, similar to Terraform Registry but specifically for OpenTofu.

#### Key Features:

- Curated collection of OpenTofu modules
- Version management
- Documentation and examples
- Community contributions
- Compatible with OpenTofu-specific features

## Finding Modules on Library.tf

1. Visit <https://library.tf>
2. Search by provider (AWS, Azure, GCP, etc.)
3. Filter by category (networking, compute, security, etc.)
4. Review module documentation and examples
5. Check version history and compatibility

## Using Library.tf Modules

### Basic Usage:

```
module "vpc" {
 source = "registry.opentofu.org/terraform-aws-modules/vpc/aws"
 version = "5.1.0"

 name = "my-vpc"
 cidr = "10.0.0.0/16"

 azs = ["us-east-1a", "us-east-1b", "us-east-1c"]
 private_subnets = ["10.0.1.0/24", "10.0.2.0/24", "10.0.3.0/24"]
 public_subnets = ["10.0.101.0/24", "10.0.102.0/24", "10.0.103.0/24"]

 enable_nat_gateway = true
 enable_vpn_gateway = false

 tags = {
 Environment = "dev"
 ManagedBy = "OpenTofu"
 }
}
```

## Evaluating Community Modules

### Quality Checklist:

- Well-documented with clear examples
- Active maintenance (recent commits)
- Proper versioning scheme
- Comprehensive input variables
- Useful output values
- Tests or validation examples
- Clear license
- Community usage and stars

### Security Considerations:

- Review source code before use
- Pin to specific versions
- Check for sensitive data handling
- Verify provider requirements
- Test in non-production first

# Hands-On Lab: Build and Reuse a VPC Module with Versioning

In this lab, you will:

1. Create a reusable VPC module from scratch
2. Use the module locally in a root configuration
3. Version the module using Git tags
4. Publish the module to GitHub
5. Consume the module with version constraints
6. Iterate on the module and release a new version

## Prerequisites:

- OpenTofu installed (1.6.0 or later)
- AWS CLI configured with credentials
- Git installed and configured
- GitHub account (for publishing)
- Text editor or IDE

**What You'll Build:** A production-ready VPC module that creates:

- VPC with customizable CIDR
- Public and private subnets across multiple AZs
- Internet Gateway
- NAT Gateway (optional)
- Route tables and associations
- Proper tagging and naming

# **OpenTofu State Management & Remote Backends**

# Learning Objectives

By the end of this module, you will be able to:

- Understand the purpose and structure of OpenTofu state
- Compare local and remote state storage approaches
- Configure remote backends (S3, GCS, Azure Blob Storage)
- Implement state locking using DynamoDB
- Import existing infrastructure into state
- Manipulate state safely using state commands
- Migrate from local to remote state storage

## Prerequisites

- Basic understanding of OpenTofu resources and providers
- AWS account with appropriate permissions
- OpenTofu CLI installed (v1.6 or later)
- AWS CLI configured with credentials

## 1. Understanding OpenTofu State

### What is State?

OpenTofu state is a JSON file that maps your configuration to real-world resources. It serves as the "source of truth" for your infrastructure.

### Key purposes of state:

- **Resource tracking:** Maintains metadata about resources managed by OpenTofu
- **Performance optimization:** Caches resource attributes to avoid constant API calls
- **Dependency mapping:** Tracks relationships between resources
- **Drift detection:** Enables comparison between desired and actual state

## State File Structure

A typical state file contains:

```
{
 "version": 4,
 "terraform_version": "1.6.0",
 "serial": 5,
 "lineage": "unique-uuid",
 "outputs": {},
 "resources": [
 {
 "mode": "managed",
 "type": "aws_instance",
 "name": "web",
 "provider": "provider[\"registry.opentofu.org/hashicorp/aws\"]",
 "instances": [...]
 }
]
}
```

### Important fields:

- `version`: State file format version
- `serial`: Increments with each state modification
- `lineage`: Unique identifier for the state history
- `resources`: Array of all managed resources

### The Dangers of State

State files contain sensitive information:

- Resource IDs and ARNs
- IP addresses and DNS names
- Database passwords (if stored in plaintext)
- API keys and secrets

**Never commit state files to version control!**

## 2. Local vs Remote State

### Local State

#### How it works:

- State stored in `terraform.tfstate` file in working directory
- Simple, no additional configuration needed
- Suitable for personal projects and learning

#### Advantages:

- Simple setup, no dependencies
- Fast read/write operations
- No network requirements
- Easy to inspect and debug

#### Disadvantages:

- No collaboration support (single user only)
- No state locking (risk of concurrent modifications)
- Risk of accidental deletion or corruption
- Difficult to share across teams
- No encryption at rest by default
- State contains sensitive data on local disk

**When to use local state:**

- Personal learning projects
- Proof of concepts
- Temporary test environments
- Single-user development scenarios

## **Remote State**

### **How it works:**

- State stored in a remote location (S3, GCS, Azure Blob, etc.)
- Multiple team members can access the same state
- Supports state locking to prevent concurrent modifications
- Often includes encryption and versioning

### **Advantages:**

- Team collaboration enabled
- State locking prevents conflicts
- Encryption at rest and in transit
- Versioning and backup capabilities
- Centralized state management
- Audit trail of state changes
- Better security controls

### **Disadvantages:**

- Additional infrastructure required
- Network dependency
- Slightly slower operations
- More complex setup
- Potential costs for storage and API calls

### **When to use remote state:**

- Team environments
- Production infrastructure
- CI/CD pipelines
- Any shared infrastructure

### 3. Remote Backend Options

#### S3 Backend (AWS)

```
terraform {
 backend "s3" {
 bucket = "my-tofu-state"
 key = "prod/terraform.tfstate"
 region = "us-east-1"
 encrypt = true
 dynamodb_table = "tofu-state-lock"

 # Optional: Use specific profile
 profile = "production"

 # Optional: Server-side encryption
 kms_key_id = "arn:aws:kms:us-east-1:123456789012:key/12345678-1234-1234-123
 }
}
```

#### Features:

- Encryption at rest using AES-256 or KMS
- Versioning support for state history
- State locking via DynamoDB
- Cost-effective for most use cases
- Integrates with AWS IAM for access control

#### Requirements:

- S3 bucket with versioning enabled
- DynamoDB table for state locking
- IAM permissions for S3 and DynamoDB

## GCS Backend (Google Cloud)

```
terraform {
 backend "gcs" {
 bucket = "my-tofu-state"
 prefix = "prod"

 # Optional: Encryption
 encryption_key = "base64-encoded-encryption-key"

 # Optional: Specific credentials
 credentials = "path/to/service-account-key.json"
 }
}
```

### Features:

- Automatic encryption at rest
- Object versioning
- Built-in state locking (no additional service needed)
- IAM integration
- Audit logging via Cloud Audit Logs

### Requirements:

- GCS bucket
- Service account with appropriate permissions
- Google Cloud SDK configured

## Azure Blob Storage Backend

```
terraform {
 backend "azurerm" {
 resource_group_name = "tfstate-rg"
 storage_account_name = "tfstateteststorage"
 container_name = "tfstate"
 key = "prod.terraform.tfstate"

 # Optional: Use specific subscription
 subscription_id = "00000000-0000-0000-0000-000000000000"

 # Optional: Use managed identity
 use_msi = true
 }
}
```

### Features:

- Encryption at rest enabled by default
- Built-in state locking via blob leases
- Versioning support
- Azure RBAC integration
- Soft delete capability

### Requirements:

- Azure Storage Account
- Storage container
- Appropriate Azure AD permissions

## Other Backend Options

### HTTP Backend:

- Custom REST API endpoints
- Useful for custom solutions
- Requires implementing the backend protocol

### Consul Backend:

- HashiCorp Consul key/value store
- Built-in locking
- Good for HashiCorp stack integration

### Terraform Cloud/Enterprise:

- Managed remote state
- Built-in VCS integration
- Policy enforcement
- Cost estimation

## 4. State Locking with DynamoDB

### Why State Locking Matters

**The Problem:** Two team members run `tofu apply` simultaneously:

1. Both read the current state
2. Both make changes
3. Both try to write state
4. One overwrites the other's changes
5. Infrastructure drift and lost changes occur

**The Solution:** State locking ensures only one operation can modify state at a time.

## How DynamoDB Locking Works

1. Before modifying state, OpenTofu attempts to acquire a lock
2. Lock information is written to DynamoDB table
3. If lock exists, operation waits or fails
4. After state update, lock is released
5. Lock includes metadata: who, when, operation ID

### Key points:

- Table name referenced in backend configuration
- `LockID` is the primary key (String type)
- `PAY_PER_REQUEST` is cost-effective for most use cases
- No additional attributes needed

## DynamoDB Table Setup

**Required table structure:**

```
resource "aws_dynamodb_table" "tofu_locks" {
 name = "tofu-state-lock"
 billing_mode = "PAY_PER_REQUEST"
 hash_key = "LockID"

 attribute {
 name = "LockID"
 type = "S"
 }

 tags = {
 Name = "OpenTofu State Lock Table"
 Environment = "Production"
 }
}
```

## Lock Behavior

### Automatic locking:

- `tofu apply` - Locks state during plan and apply
- `tofu destroy` - Locks during destroy operation
- `tofu refresh` - Locks during refresh
- `tofu import` - Locks during import

### Force unlock:

```
If a lock gets stuck (use with extreme caution)
tofu force-unlock <LOCK_ID>
```

**Warning:** Only force-unlock if you're certain no other process is running!

## **Handling Lock Conflicts**

## Common scenarios:

### Scenario 1: Legitimate concurrent operation

Error: Error acquiring the state lock

Lock Info:

```
ID: abc123-def456-ghi789
Path: my-bucket/terraform.tfstate
Operation: OperationTypeApply
Who: john@example.com
Version: 1.6.0
Created: 2024-01-15 10:30:45 UTC
```

### Scenario 2: Stale lock (process crashed)

bash

```
Verify no process is actually running
Then force unlock
tofu force-unlock abc123-def456-ghi789
```

**Resolution:** Wait for the other operation to complete.

## 5. State Manipulation Commands

### Viewing State

#### List all resources:

```
bash
```

```
tofu state list
```

#### Example output:

```
aws_vpc.main
aws_subnet.public[0]
aws_subnet.public[1]
aws_instance.web
```

#### Show specific resource:

```
bash
```

```
tofu state show aws_instance.web
```

#### Show all state details:

```
bash
```

```
tofu show
```

## Importing Existing Resources

**Use case:** Bring existing infrastructure under OpenTofu management.

### Basic import:

---

bash

```
tofu import aws_instance.web i-1234567890abcdef0
```

## Import with configuration:

1. Create the resource block in your configuration:

```
hcl

resource "aws_instance" "web" {
 # Configuration will be filled in after import
 ami = "ami-12345678"
 instance_type = "t3.micro"
}
```

2. Import the existing resource:

```
bash

tofu import aws_instance.web i-1234567890abcdef0
```

3. Run plan to see any drift:

```
bash

tofu plan
```

4. Update configuration to match actual state (eliminate drift)

### **Import with modules:**

```
bash
```

```
tofu import module.vpc.aws_vpc.main vpc-12345678
```

### **Import with for\_each:**

```
bash
```

```
tofu import 'aws_subnet.public["us-east-1a"]' subnet-12345678
```

## Moving Resources in State

### Rename a resource:

```
bash
```

```
tofu state mv aws_instance.web aws_instance.webserver
```

### Move to a module:

```
bash
```

```
tofu state mv aws_instance.web module.web.aws_instance.server
```

### Move from a module:

```
bash
```

```
tofu state mv module.web.aws_instance.server aws_instance.web
```

### Move between indexed resources:

```
bash
```

```
tofu state mv 'aws_subnet.public[0]' 'aws_subnet.public[1]'
```

**Common use cases:**

- Refactoring module structure
- Renaming resources without recreating them
- Reorganizing resource hierarchies

## Removing Resources from State

### Remove without destroying:

```
bash
```

```
tofu state rm aws_instance.old_server
```

### What happens:

- Resource removed from state
- Physical resource remains unchanged
- OpenTofu will no longer manage it
- Next apply won't destroy it

### Use cases:

- Handing off management to another tool
- Removing resources before deletion
- Cleaning up after failed imports

## Replacing Resources

### Mark resource for replacement:

```
bash
```

```
tofu apply -replace=aws_instance.web
```

### What happens:

- Resource is destroyed and recreated
- Even if configuration hasn't changed
- Useful for troubleshooting

## Pulling and Pushing State

### **Pull current state:**

---

bash

```
tofu state pull > terraform.tfstate.backup
```

---

### **Push modified state (DANGEROUS):**

---

bash

```
tofu state push terraform.tfstate.modified
```

---

**Warning:** Manual state push can corrupt state. Only use in recovery scenarios!

## State Command Best Practices

### 1. Always backup before manipulation:

---

bash

```
tofu state pull > backup-$(date +%Y%m%d-%H%M%S).tfstate
```

---

### 2. Use `-dry-run` when available:

---

bash

```
tofu state mv -dry-run aws_instance.old aws_instance.new
```

---

### 3. Verify changes:

---

bash

```
tofu plan
```

---

### 4. Use state locking: Ensure backend has locking enabled

### 5. Document changes: Keep notes on why state was manipulated

## **6. State Migration Strategies**

### **Migration Planning**

#### **Questions to answer:**

- Where is state currently stored?
- Where should it be stored?
- Is state locking configured?
- Are there multiple state files to migrate?
- Do team members need access?

# Migration Process

## Step 1: Backup current state

bash

```
cp terraform.tfstate terraform.tfstate.backup
```

## Step 2: Create remote backend resources

- S3 bucket
- DynamoDB table
- IAM permissions

## Step 3: Add backend configuration

hcl

```
terraform {
 backend "s3" {
 bucket = "my-tofu-state"
 key = "terraform.tfstate"
 region = "us-east-1"
 encrypt = true
 dynamodb_table = "tofu-locks"
 }
}
```

## Step 4: Initialize with migration

bash

```
tofu init -migrate-state
```

## Step 5: Verify migration

bash

```
tofu plan
Should show no changes if migration successful
```

## Step 6: Clean up local state (after verification)

bash

```
Move to archive, don't delete immediately
mkdir -p old-state
mv terraform.tfstate* old-state/
```

## Partial Backend Configuration

**Use case:** Different backends per environment without hardcoding values.

### **backend.hcl:**

```
hcl

bucket = "prod-tofu-state"
key = "prod/terraform.tfstate"
region = "us-east-1"
encrypt = true
dynamodb_table = "prod-tofu-locks"
```

### **Initialize with partial config:**

```
bash

tofu init -backend-config=backend.hcl
```

### **Benefits:**

- Keep sensitive values out of version control
- Use different configs per environment
- CI/CD friendly approach

## Access Control

### S3 Bucket Policy:

---

hcl

```
{
 "Version": "2012-10-17",
 "Statement": [
 {
 "Effect": "Allow",
 "Principal": {
 "AWS": "arn:aws:iam::123456789012:role/Tofu-Role"
 },
 "Action": [
 "s3:GetObject",
 "s3:PutObject",
 "s3:DeleteObject"
],
 "Resource": "arn:aws:s3:::my-tofu-state/*"
 },
 {
 "Effect": "Allow",
 "Principal": {
 "AWS": "arn:aws:iam::123456789012:role/Tofu-Role"
 },
 "Action": "s3>ListBucket",
 "Resource": "arn:aws:s3:::my-tofu-state"
 }
]
}
```

## 7. Security Best Practices

### Encryption

#### S3 Backend:

- Enable default encryption on bucket
- Use KMS for enhanced security
- Enable bucket versioning

#### Secrets in State:

- Avoid storing secrets in plaintext
- Use AWS Secrets Manager or Vault
- Reference secrets, don't embed them

## DynamoDB Permissions

### Required permissions:

```
json
{
 "Version": "2012-10-17",
 "Statement": [
 {
 "Effect": "Allow",
 "Action": [
 "dynamodb:GetItem",
 "dynamodb:PutItem",
 "dynamodb:DeleteItem"
],
 "Resource": "arn:aws:dynamodb:*:*:table/tofu-state-lock"
 }
]
}
```

### Audit and Monitoring

#### Enable CloudTrail logging:

- Track who accessed state
- Monitor state modifications
- Alert on suspicious activity

#### S3 Access Logging:

- Log all bucket access
- Review access patterns
- Detect unauthorized access

## 8. Common Patterns and Anti-Patterns

### Good Patterns

#### **Environment Separation:**

```
s3://my-state-bucket/dev/terraform.tfstate
s3://my-state-bucket/staging/terraform.tfstate
s3://my-state-bucket/prod/terraform.tfstate
```

#### **Workspace-Based Organization:**

- Use OpenTofu workspaces for similar environments
- Separate buckets for truly different systems

#### **State File Naming:**

- Use descriptive key names
- Include environment information
- Use consistent structure

## Anti-Patterns

### Don't:

- Store state in version control
- Share state files via email or chat
- Manually edit state files
- Use same state for unrelated infrastructure
- Skip state backups
- Ignore state locking

### Don't store secrets directly in variables:

```
hcl

Bad
variable "db_password" {
 default = "SuperSecret123!" # Will end up in state!
}

Good
variable "db_password" {
 sensitive = true
 # Pass via environment variable or secrets manager
}
```

## Summary

### Key Takeaways:

1. **State is critical:** It's the mapping between your config and real resources
2. **Remote state for teams:** Always use remote state in team environments
3. **Enable locking:** Prevent concurrent modifications with DynamoDB (S3) or built-in locking (GCS, Azure)
4. **Import carefully:** Bring existing resources under management systematically
5. **Manipulate with caution:** Always backup before using state commands
6. **Secure your state:** Encrypt at rest, control access, audit changes
7. **Plan migrations:** Test in non-production first, always backup

**LAB: Migrate to Remote S3 Backend and Manage Concurrent Applies**

## Lab Overview

**Scenario:** You're a DevOps engineer at a growing startup. Your infrastructure was initially managed by a single developer using local state. Now that the team is expanding, you need to migrate to a remote backend to enable team collaboration and prevent state conflicts.

### What You'll Do:

Deploy infrastructure using local state

Create S3 bucket and DynamoDB table for remote state

Migrate from local to remote state

Test state locking with concurrent operations

Practice state manipulation commands (import, move, remove)

Handle lock conflicts and recovery scenarios

### Learning Outcomes:

Configure and migrate to S3 remote backend

Implement state locking with DynamoDB

Safely manipulate state

Troubleshoot common state issues

## Prerequisites

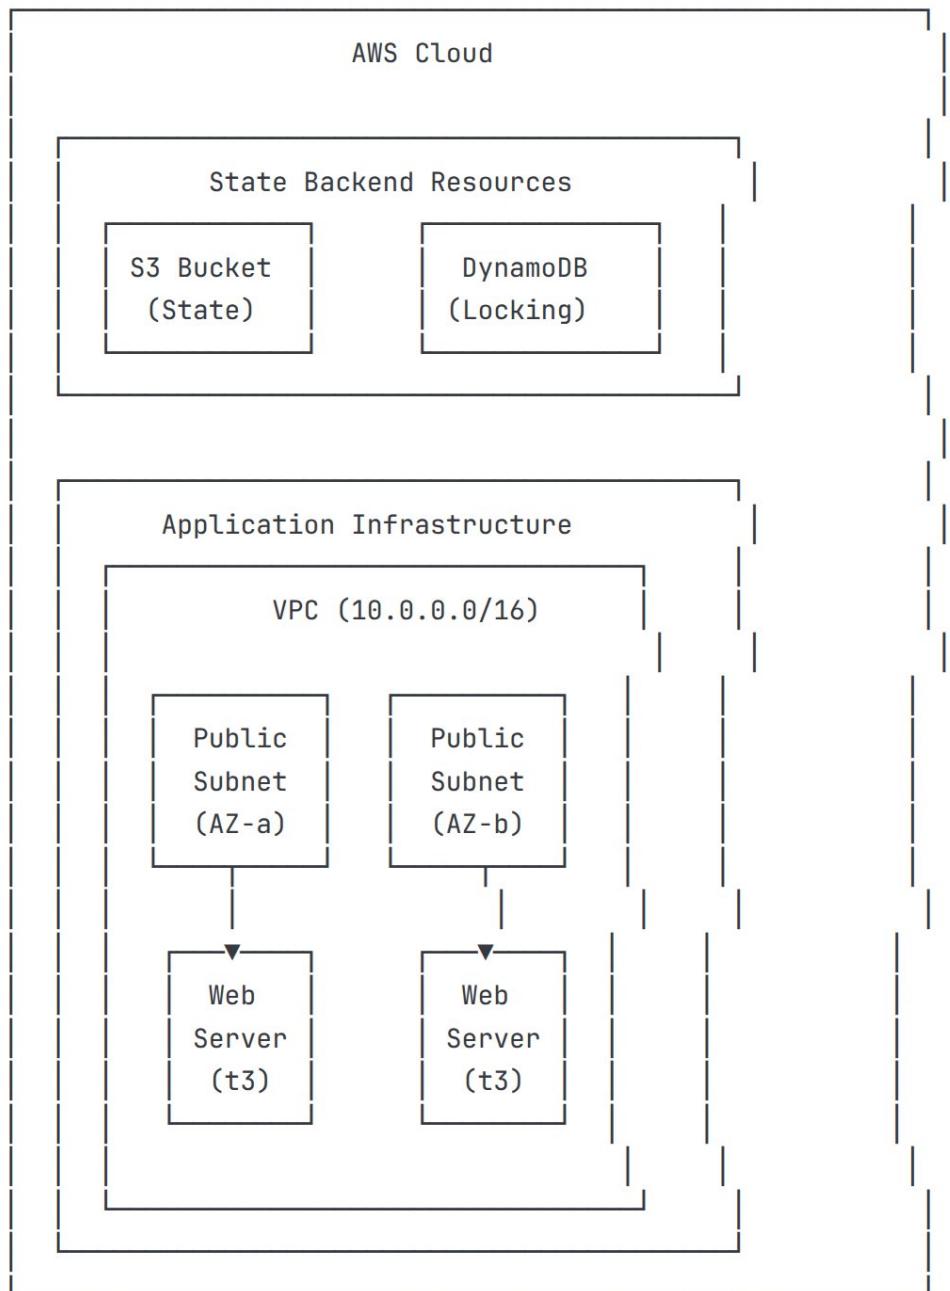
### Required Tools

- OpenTofu CLI (v1.6+): `tofu --version`
- AWS CLI configured: `aws sts get-caller-identity`
- Text editor
- Terminal with multiple tabs/windows (for concurrent testing)

### AWS Permissions Required

- S3: CreateBucket, PutObject, GetObject, ListBucket
- DynamoDB: CreateTable, PutItem, GetItem, DeleteItem
- EC2: CreateVpc, CreateSubnet, RunInstances, DescribeInstances
- IAM: PassRole (if using instance profiles)

## Lab Architecture



# **OpenTofu Provisioners & Connections**

# Module Overview

## **Learning Objectives:**

- Understand what provisioners are and when to use them
- Configure local-exec, remote-exec, and file provisioners
- Set up SSH connections for remote execution
- Implement error handling strategies
- Deploy and configure EC2 instances post-deployment

## **Prerequisites:**

- Basic OpenTofu knowledge (providers, resources, variables)
- Familiarity with AWS EC2
- Basic Linux command-line skills
- Understanding of SSH concepts

## 1. Understanding Provisioners

### What Are Provisioners?

Provisioners are OpenTofu's mechanism for executing actions on local or remote machines as part of resource creation or destruction. They allow you to:

- Run configuration management scripts
- Initialize software on newly created instances
- Execute cleanup tasks during resource destruction
- Transfer files to remote systems
- Trigger external processes

**Provisioners** are like "after-market scripts."

They are commands that execute either on your local machine or on the newly created server immediately after the infrastructure is built.

## 1. The Analogy: Building a House

- **OpenTofu (The Builder):** Builds the house. It puts up the walls, installs the roof, and connects the electricity. The resource is "created."
- **Provisioner (The Movers):** Once the house is finished, the movers come in and arrange the furniture, hang paintings, and stock the fridge.
  - *Without provisioners*, you have an empty house (a blank server).
  - *With provisioners*, you have a house ready to live in (a server with software installed).

## When to Use Provisioners

 **Important:** Provisioners are a "last resort" in OpenTofu. They should be used when no better alternative exists.

### Good Use Cases:

- Bootstrapping initial configuration management (Ansible, Chef, Puppet)
- One-time setup tasks that can't be handled by cloud-init
- Triggering external systems that lack API integration
- Development and testing scenarios

### Better Alternatives:

- **cloud-init / user\_data** - For instance initialization
- **Configuration management tools** - For ongoing configuration
- **Custom AMIs** - For pre-configured images
- **Container orchestration** - For application deployment
- **AWS Systems Manager** - For instance configuration

### Provisioner Lifecycle

Provisioners run:

- **Creation-time:** After resource is created (default)
- **Destruction-time:** Before resource is destroyed (when = destroy)

**Critical Limitation:** Provisioners only run during resource creation or destruction, NOT during updates. This makes them unsuitable for ongoing configuration management.

#### 4. When Should You Use Them? (The "Last Resort" Rule)

In modern DevOps, **you should avoid Provisioners if possible.**

##### Why?

- **They are "Dumb":** OpenTofu doesn't know what happens inside the script. If the script fails halfway, Tofu might think the server is broken ("tainted") and try to delete it next time.
- **No State:** Tofu tracks the server, but it doesn't track the *software* inside. If you change the script, Tofu won't know to update the server.

## 2. Types of Provisioners

### 2.1 local-exec Provisioner

Executes commands on the **machine running OpenTofu** (your workstation or CI/CD server).

#### Syntax:

---

```
hcl

provisioner "local-exec" {
 command = "echo ${self.private_ip} >> private_ips.txt"

 # Optional parameters
 working_dir = "/tmp"
 interpreter = ["/bin/bash", "-c"]
 environment = {
 FOO = "bar"
 }
}
```

### **Use Cases:**

- Recording information locally (IP addresses, DNS names)
- Triggering local scripts or tools
- Updating local configuration files
- Calling external APIs or webhooks
- Running local validation checks

```
Update local DNS records
provisioner "local-exec" {
 command = "aws route53 change-resource-record-sets --cli-input-json file://dns-update.json"
}

Save instance information
provisioner "local-exec" {
 command = "echo 'Instance ${self.id} created at ${self.public_ip}' >> deployment.log"
}

Trigger external monitoring
provisioner "local-exec" {
 command = "curl -X POST https://monitoring.example.com/register -d 'host=${self.private_ip}'"
}
```

### Interpreter Options:

- Default: `["/bin/sh", "-c"]` on Unix, `["cmd", "/C"]` on Windows
- Can specify custom: `interpreter = ["/bin/bash", "-c"]`
- PowerShell: `interpreter = ["PowerShell", "-Command"]`

## 2.2 remote-exec Provisioner

Executes commands on the **remote resource** after it's created.

### Syntax:

```
hcl

provisioner "remote-exec" {
 inline = [
 "sudo apt-get update",
 "sudo apt-get install -y nginx",
 "sudo systemctl start nginx"
]

 # Or use a script file
 script = "scripts/setup.sh"

 # Or use multiple scripts
 scripts = ["scripts/install.sh", "scripts/configure.sh"]
}
```

### Three Execution Methods:

#### 1. **inline** - List of commands (most common)

```
hcl

inline = [
 "sudo yum update -y",
 "sudo yum install -y docker",
 "sudo systemctl enable docker",
 "sudo systemctl start docker"
]
```

#### Use Cases:

- Installing software packages
- Starting services
- Running configuration scripts
- Waiting for services to be ready
- Executing deployment commands

#### 2. **script** - Single local script file to execute

```
hcl

script = "${path.module}/scripts/bootstrap.sh"
```

#### 3. **scripts** - Multiple script files in order

```
hcl

scripts = [
 "${path.module}/scripts/install-deps.sh",
 "${path.module}/scripts/configure-app.sh",
 "${path.module}/scripts/start-services.sh"
]
```

## 2.3 file Provisioner

Copies files or directories from local machine to remote resource.

### Syntax:

```
hcl

provisioner "file" {
 source = "conf/myapp.conf"
 destination = "/etc/myapp.conf"
}

Copy entire directory
provisioner "file" {
 source = "configs/" # Trailing slash important!
 destination = "/etc/myapp"
}

Use content instead of file
provisioner "file" {
 content = templatefile("${path.module}/config.tpl", {
 server_port = var.port
 })
 destination = "/tmp/config.txt"
}
```

### Parameters:

- **source** - Local file or directory path
- **destination** - Remote path (must be writable by connection user)
- **content** - Direct content (alternative to source)

### Directory Behavior:

- `source = "foo"` → Copies directory `foo` into destination
- `source = "foo/"` → Copies contents of `foo` into destination (note trailing slash)

### Use Cases:

- Deploying configuration files
- Transferring application code
- Uploading SSL certificates
- Copying installation packages
- Deploying scripts before execution

## Typical Pattern:

```
hcl

First, copy the script
provisioner "file" {
 source = "${path.module}/scripts/setup.sh"
 destination = "/tmp/setup.sh"
}

Then, execute it
provisioner "remote-exec" {
 inline = [
 "chmod +x /tmp/setup.sh",
 "sudo /tmp/setup.sh"
]
}
```

### 3. Connection Blocks

Connection blocks define how OpenTofu connects to remote resources for `remote-exec` and `file` provisioners.

#### SSH Connections (Linux/Unix)

```
hcl

connection {
 type = "ssh"
 host = self.public_ip
 user = "ec2-user"
 private_key = file("~/ssh/id_rsa")
 timeout = "5m"
}
```

## SSH Connection Parameters:

| Parameter                   | Description              | Default | Example                            |
|-----------------------------|--------------------------|---------|------------------------------------|
| <code>type</code>           | Connection type          | "ssh"   | "ssh"                              |
| <code>host</code>           | Remote host              | -       | <code>self.public_ip</code>        |
| <code>port</code>           | SSH port                 | 22      | 2222                               |
| <code>user</code>           | SSH username             | "root"  | "ec2-user"                         |
| <code>password</code>       | SSH password             | -       | "secret123"                        |
| <code>private_key</code>    | SSH private key content  | -       | <code>file("~/ssh/key.pem")</code> |
| <code>certificate</code>    | SSH certificate          | -       | <code>file("cert.pub")</code>      |
| <code>agent</code>          | Use SSH agent            | true    | false                              |
| <code>agent_identity</code> | Specific key from agent  | -       | "key-name"                         |
| <code>host_key</code>       | Expected host public key | -       | -                                  |
| <code>timeout</code>        | Connection timeout       | "5m"    | "10m"                              |
| <code>script_path</code>    | Remote path for scripts  | -       | "/tmp/script.sh"                   |

## WinRM Connections (Windows)

```
hcl
connection {
 type = "winrm"
 host = self.public_ip
 user = "Administrator"
 password = var.admin_password
 https = true
 insecure = true
 timeout = "10m"
}
```

## WinRM Connection Parameters:

| Parameter             | Description                 | Default                   |
|-----------------------|-----------------------------|---------------------------|
| <code>type</code>     | Connection type             | "winrm"                   |
| <code>host</code>     | Remote host                 | -                         |
| <code>port</code>     | WinRM port                  | 5985 (HTTP), 5986 (HTTPS) |
| <code>user</code>     | Username                    | -                         |
| <code>password</code> | Password                    | -                         |
| <code>https</code>    | Use HTTPS                   | false                     |
| <code>insecure</code> | Skip certificate validation | false                     |
| <code>use_ntlm</code> | Use NTLM auth               | false                     |
| <code>timeout</code>  | Connection timeout          | "5m"                      |

## Connection Block Placement

### Option 1: Inside Provisioner (Specific)

```
hcl

provisioner "remote-exec" {
 connection {
 type = "ssh"
 host = self.public_ip
 user = "ubuntu"
 private_key = file(var.private_key_path)
 }
 inline = ["echo 'Connected!'"]
}
```

### Option 2: Inside Resource (Shared by all provisioners)

```
hcl

resource "aws_instance" "web" {
 # ... instance configuration ...

 connection {
 type = "ssh"
 host = self.public_ip
 user = "ec2-user"
 private_key = file(var.private_key_path)
 }

 provisioner "file" {
 source = "app.conf"
 destination = "/tmp/app.conf"
 }

 provisioner "remote-exec" {
 inline = ["sudo mv /tmp/app.conf /etc/app.conf"]
 }
}
```

## Using self Object

Inside provisioners, use `self` to reference the parent resource's attributes:

```
hcl

connection {
 host = self.public_ip # Instance's public IP
 user = self.tags["SSHUser"] # Tag value from instance
}

provisioner "local-exec" {
 command = "echo ${self.id} >> instances.txt" # Instance ID
}
```

## 4. Error Handling

### on\_failure Behavior

Controls what happens when a provisioner fails.

#### Syntax:

---

```
hcl

provisioner "remote-exec" {
 inline = ["exit 1"] # This will fail

 on_failure = continue # or "fail"
}
```

---

#### Options:

| Value                 | Behavior                                           |
|-----------------------|----------------------------------------------------|
| <code>fail</code>     | Stop and mark resource as tainted (default)        |
| <code>continue</code> | Log error but continue, resource marked as created |

## Example Use Cases:

```
hcl

Critical provisioner - must succeed
provisioner "remote-exec" {
 inline = [
 "sudo systemctl start critical-service"
]
 on_failure = fail # Default behavior
}

Optional provisioner - failure OK
provisioner "remote-exec" {
 inline = [
 "sudo apt-get update",
 "sudo apt-get upgrade -y" # Might fail, but OK
]
 on_failure = continue
}

Cleanup task - don't fail if already clean
provisioner "remote-exec" {
 inline = [
 "rm -f /tmp/setup.sh || true"
]
 on_failure = continue
}
```

## **when** Parameter

Controls when provisioner runs.

### **Syntax:**

---

hcl

```
provisioner "local-exec" {
 when = destroy
 command = "echo 'Instance ${self.id} being destroyed' >> cleanup.log"
}
```

---

### **Options:**

| Value   | Runs                               |
|---------|------------------------------------|
| create  | During resource creation (default) |
| destroy | During resource destruction        |

---

## Destruction-time Example:

---

```
hcl

resource "aws_instance" "app" {
 ami = data.aws_ami.amazon_linux.id
 instance_type = "t3.micro"

 # Runs during creation
 provisioner "remote-exec" {
 inline = [
 "sudo systemctl start myapp"
]
 }

 # Runs during destruction
 provisioner "remote-exec" {
 when = destroy
 inline = [
 "sudo systemctl stop myapp",
 "sudo /opt/cleanup.sh"
]
 }

 # Local cleanup during destruction
 provisioner "local-exec" {
 when = destroy
 command = "curl -X DELETE https://api.example.com/servers/${self.id}"
 }
}
```

### Important Notes:

- Destroy provisioners run **before** the resource is destroyed
- If destroy provisioner fails, resource is NOT destroyed
- `self` references still work in destroy provisioners
- Connection must still be valid (resource must exist and be accessible)

## Summary

### Key Takeaways:

1. Provisioners are a last resort - prefer cloud-init, AMIs, or config management
2. Three types: local-exec (local), remote-exec (remote), file (transfer)
3. Connection blocks define SSH/WinRM access for remote provisioners
4. Use `on_failure` and `when` for error handling
5. Only run on create/destroy, not on updates
6. Keep provisioners simple and idempotent

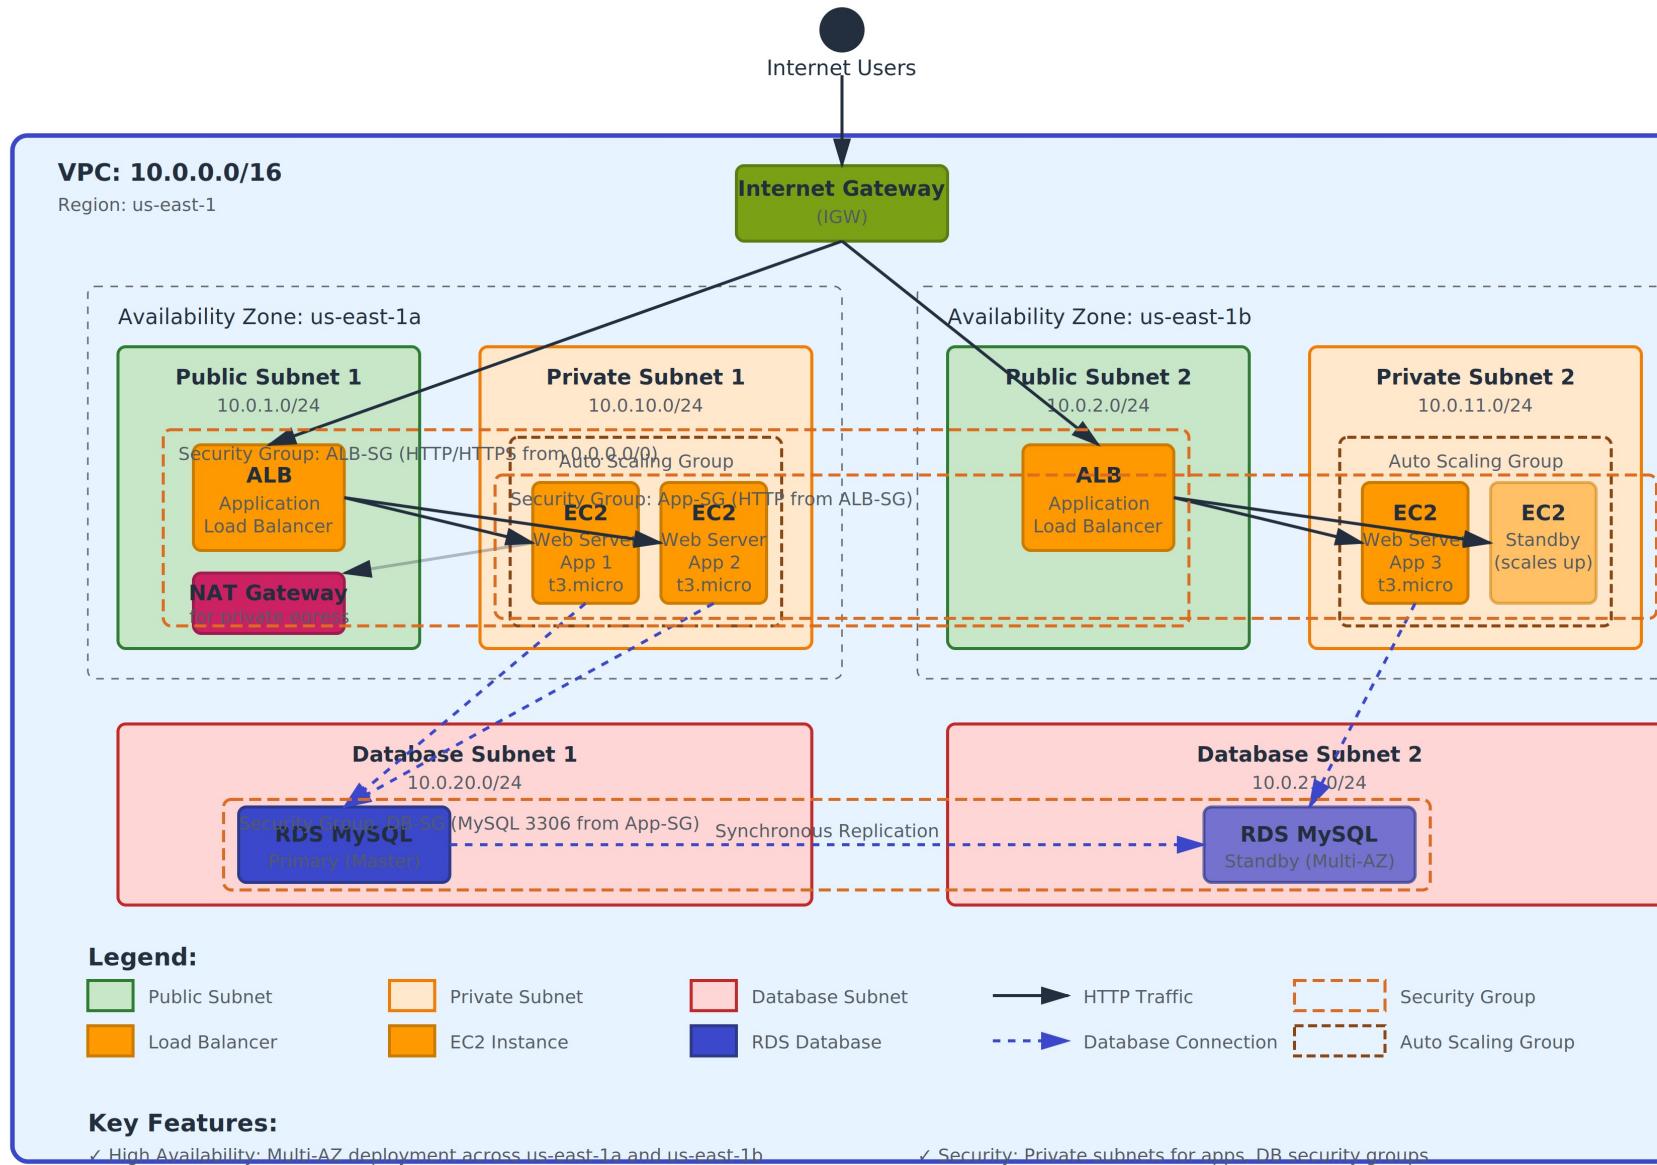
Lab 15 - Provisioners

Lab - Complex Infrastructure with  
Dependencies , Autoscaling, ALB

lesson6\_complex\_infra\_w\_depe\_lb

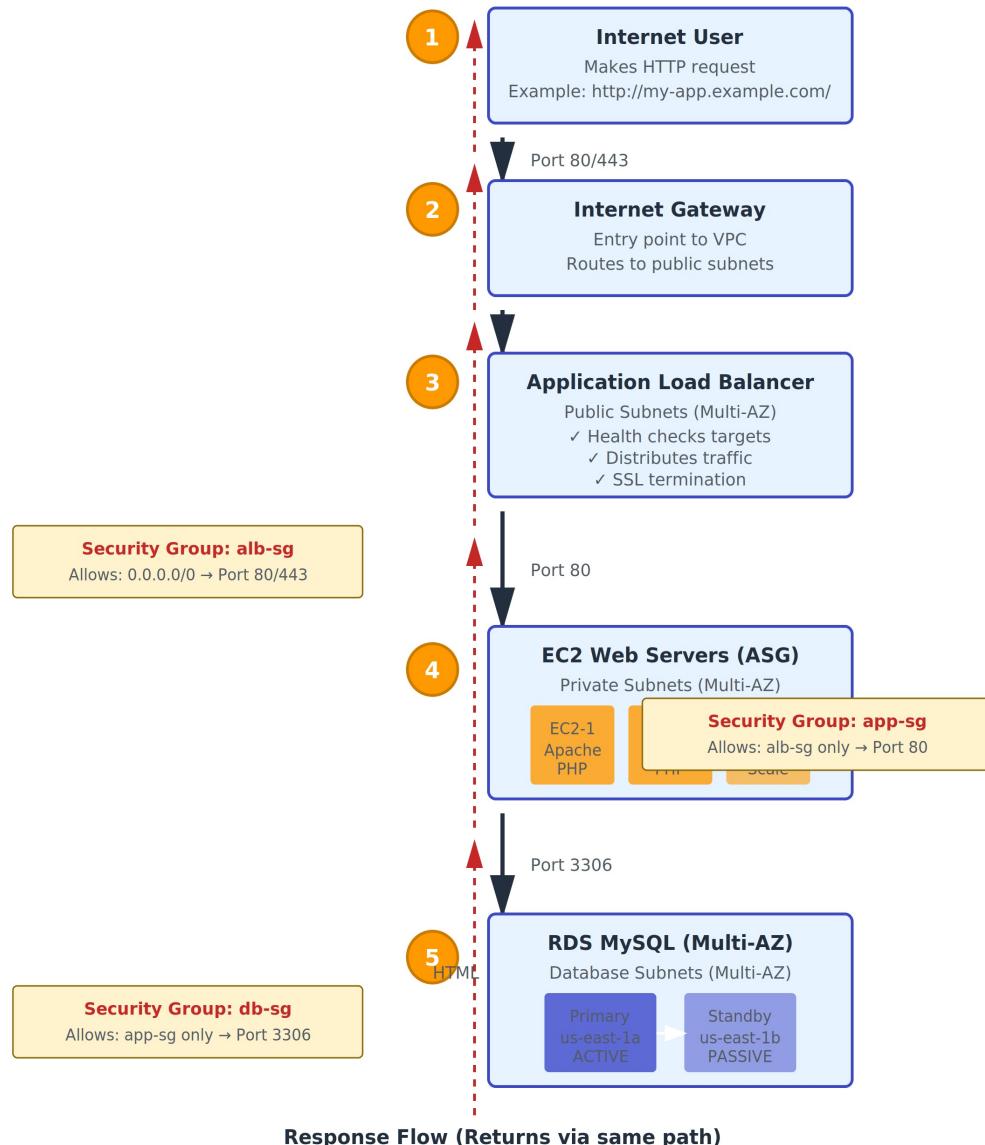
# Use Case 5: 3-Tier Web Application with Auto-Scaling

AWS Architecture - Load Balancer, Auto Scaling Group, and RDS



## Request Flow: User to Database

Complete traffic path through 3-tier architecture



### Key Security Points:

- Each tier has its own security group
- Security groups reference each other (not IP ranges)
- Database has NO internet access
- Application servers in private subnets
- Only ALB is publicly accessible

### High Availability:

- All tiers span 2 Availability Zones
- ALB distributes across healthy targets
- Auto Scaling replaces failed instances
- RDS automatic failover (Multi-AZ)
- No single point of failure

# Hands-On Lab: EC2 Post-Deployment Configuration with Provisioners

## Lab Overview

**Objective:** Deploy an EC2 instance and use provisioners to:

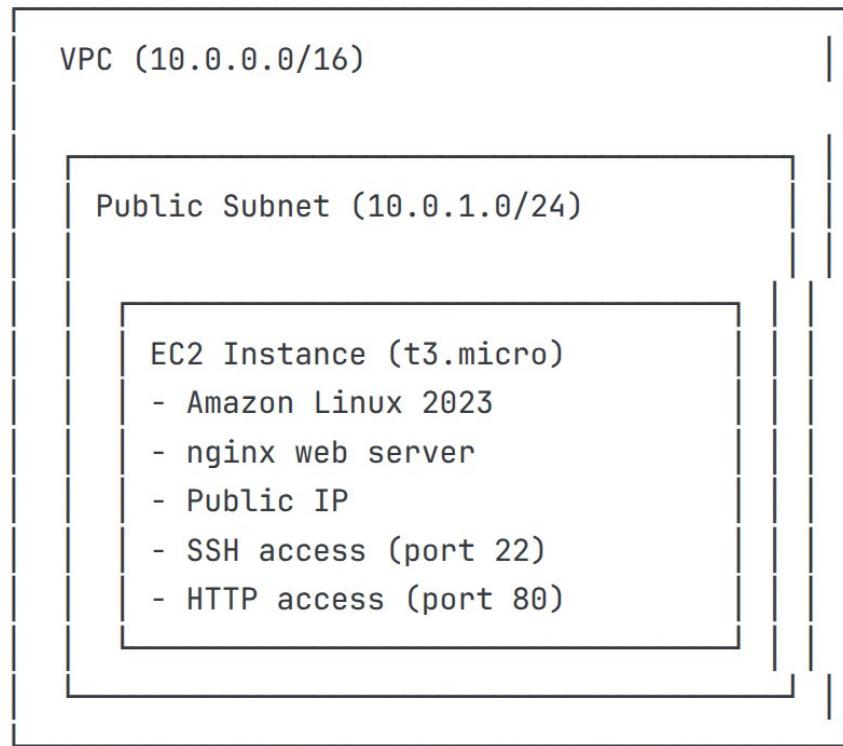
1. Install and configure a web server (nginx)
2. Deploy a custom web page
3. Configure application settings
4. Test connectivity and verify configuration

**Scenario:** You're deploying a web application server that needs post-deployment configuration. The application requires nginx installation, custom configuration files, and a welcome page with dynamic content.

### What You'll Learn:

- Configure all three provisioner types
- Set up SSH connections
- Handle provisioner errors
- Transfer files to remote instances
- Execute remote configuration commands
- Verify provisioner success

## Lab Architecture



## Prerequisites

Before starting this lab, ensure you have:

- AWS CLI configured with valid credentials
- OpenTofu installed (v1.6.0+)
- SSH key pair available (or script will create one)
- Permissions to create VPC, EC2, and security group resources
- Network connectivity to AWS

LAB - Session(5) -EC2 Post-Deployment Configuration with Provisioners

Provisioner Flow:

1. file: Copy nginx config → /tmp/nginx.conf
2. file: Copy welcome page → /tmp/index.html
3. remote-exec: Install nginx, move files, start service
4. local-exec: Record deployment info locally

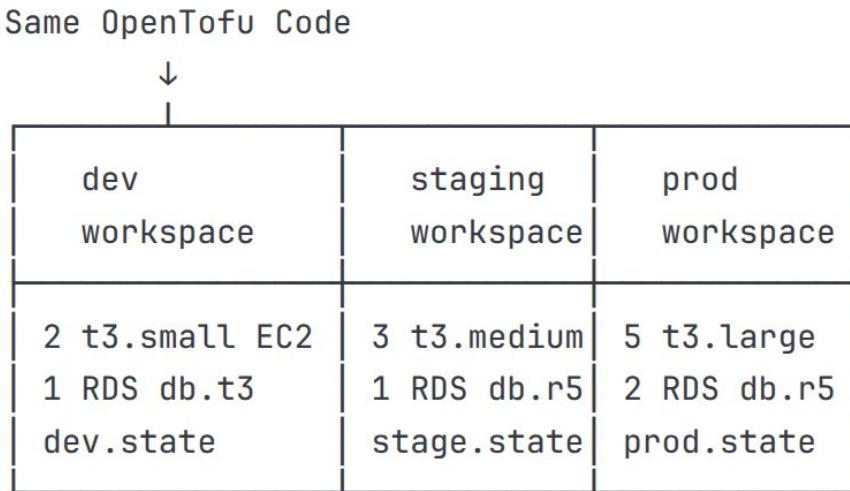
# **OpenTofu Workspaces - Concepts & Fundamentals**

## **What Are Workspaces?**

Workspaces in OpenTofu allow you to manage multiple distinct sets of infrastructure resources using the same configuration code. Each workspace has its own separate state file, enabling you to maintain different environments (dev, staging, production) from a single codebase.

## The Core Concept

Think of workspaces as parallel universes for your infrastructure:



Each workspace maintains:

- **Independent state file:** No state pollution between environments
- **Same configuration code:** DRY principle for infrastructure
- **Isolated resources:** Resources in one workspace don't affect others
- **Separate lifecycle:** Create/update/destroy independently

## How Workspaces Work Internally

### Default Workspace

Every OpenTofu project starts with a `default` workspace:

```
bash

Initialize a new project
tofu init

Check current workspace
tofu workspace show
Output: default

List all workspaces
tofu workspace list
Output: * default (asterisk shows current workspace)
```

## Creating Additional Workspaces

---

```
bash
```

```
Create new workspace
tofu workspace new dev
Output: Created and switched to workspace "dev"

Create more workspaces
tofu workspace new staging
tofu workspace new prod

List all workspaces
tofu workspace list
Output:
default
dev
staging
* prod
```

## Workspace State Storage

### Local Backend

When using local state, workspaces create separate state files:

```
.terraform/
└── terraform.tfstate.d/
 ├── dev/
 │ └── terraform.tfstate
 ├── staging/
 │ └── terraform.tfstate
 └── prod/
 └── terraform.tfstate
terraform.tfstate # default workspace state
```

### Remote Backend (S3)

With S3 backend, workspaces use key prefixes:

```
s3://my-terraform-state/
└── env:/dev/myapp/terraform.tfstate
└── env:/staging/myapp/terraform.tfstate
└── env:/prod/myapp/terraform.tfstate
└── myapp/terraform.tfstate # default workspace
```

# Workspace Interpolation

The most powerful feature is `terraform.workspace` interpolation:

## Example 1: Environment-Specific Naming

```
hcl

main.tf
resource "aws_instance" "web" {
 ami = var.ami_id
 instance_type = terraform.workspace == "prod" ? "t3.large" : "t3.small"

 tags = {
 Name = "${terraform.workspace}-web-server"
 Environment = terraform.workspace
 }
}
```

## Results across workspaces:

- **dev workspace:** Creates `dev-web-server` with `t3.small`
- **staging workspace:** Creates `staging-web-server` with `t3.small`
- **prod workspace:** Creates `prod-web-server` with `t3.large`

# Workspace Interpolation

## Example 2: Count-Based Scaling

```
hcl

variables.tf
variable "instance_count" {
 type = map(number)
 default = {
 dev = 1
 staging = 2
 prod = 5
 }
}

main.tf
resource "aws_instance" "app" {
 count = var.instance_count[terraform.workspace]
 ami = var.ami_id
 instance_type = "t3.small"

 tags = {
 Name = "${terraform.workspace}-app-${count.index + 1}"
 }
}
```

# Workspace Interpolation

## Example 3: Locals for Complex Logic

hcl

```
locals.tf
locals {
 environment_config = {
 dev = {
 instance_type = "t3.small"
 instance_count = 1
 enable_monitoring = false
 backup_retention = 1
 }
 staging = {
 instance_type = "t3.medium"
 instance_count = 2
 enable_monitoring = true
 backup_retention = 7
 }
 prod = {
 instance_type = "t3.large"
 instance_count = 5
 enable_monitoring = true
 backup_retention = 30
 }
 }
}
```

```
Get config for current workspace
env_config = local.environment_config[terraform.workspace]
}

main.tf
resource "aws_instance" "app" {
 count = local.env_config.instance_count
 ami = var.ami_id
 instance_type = local.env_config.instance_type

 monitoring = local.env_config.enable_monitoring

 tags = {
 Name = "${terraform.workspace}-app-${count.index + 1}"
 Environment = terraform.workspace
 }
}

resource "aws_db_instance" "main" {
 identifier = "${terraform.workspace}-database"
 engine = "postgres"
 instance_class = terraform.workspace == "prod" ? "db.r5.large" : "db.t3.micro"
 backup_retention_period = local.env_config.backup_retention

 tags = {
 Environment = terraform.workspace
 }
}
```

## Workspace Commands Reference

### Basic Operations

---

```
bash

Create and switch to new workspace
tofu workspace new <name>

Switch to existing workspace
tofu workspace select <name>

List all workspaces
tofu workspace list

Show current workspace
tofu workspace show

Delete a workspace (must be empty)
tofu workspace delete <name>
```

### Common Workflows

---

```
bash

Deploy to dev environment
tofu workspace select dev
tofu plan
tofu apply

Switch to staging and deploy
tofu workspace select staging
tofu apply

Check production state without switching
tofu workspace select prod
tofu state list
tofu workspace select dev # switch back
```

## When to Use Workspaces

### Good Use Cases

#### 1. Similar Infrastructure, Different Scale

- Same VPC layout, different CIDR blocks
- Same application tier, different instance counts
- Same architecture, different sizing

#### 2. Development Workflow

- Feature branches that need temporary environments
- Developer sandboxes from same template
- Testing infrastructure changes

#### 3. Single-Team Management

- One team manages all environments
- Consistent access patterns
- Shared responsibility model

#### 4. Rapid Environment Provisioning

- Spin up identical environments quickly
- Testing disaster recovery procedures
- Demo environments

## Workspace vs Other Patterns

### Workspaces:

```
my-app/
└── main.tf
└── variables.tf
└── outputs.tf
```

```
One codebase, multiple workspaces
```

### Separate Directories:

```
environments/
└── dev/
 └── main.tf
 └── variables.tf
└── staging/
 └── main.tf
 └── variables.tf
└── prod/
 └── main.tf
 └── variables.tf
```

### When to use each:

- **Workspaces:** Infrastructure is 90%+ identical
- **Directories:** Significant differences between environments

## Workspace vs Separate Repositories

### Workspaces:

```
infrastructure-repo/
└── app-platform/
 ├── main.tf
 └── ...
 ...
```

#### When to use each:

- **Workspaces:** Same team, similar infrastructure
- **Repositories:** Different teams, different compliance requirements

### Separate Repositories:

```
app-platform-dev/
app-platform-staging/
app-platform-prod/
```

# Common Patterns

## Pattern 1: Variable Files per Workspace

```
bash
```

```
Directory structure
```

```
.
├── main.tf
├── variables.tf
└── dev.tfvars
 ├── staging.tfvars
 └── prod.tfvars
```

```
bash
```

```
Usage
```

```
tofu workspace select dev
tofu apply -var-file="dev.tfvars"

tofu workspace select prod
tofu apply -var-file="prod.tfvars"
```

# Common Patterns

## Pattern 2: Workspace Validation

---

hcl

```
main.tf
locals {
 valid_workspaces = ["dev", "staging", "prod"]

 # Validate workspace name
 validate_workspace = contains(local.valid_workspaces, terraform.workspace)
 ? true
 : file("ERROR: Invalid workspace. Must be one of: ${join(", ", local.valid_workspaces)}")
}
```

# Common Patterns

## Pattern 3: Conditional Resources

```
hcl

Only create in production
resource "aws_cloudwatch_dashboard" "main" {
 count = terraform.workspace == "prod" ? 1 : 0
 dashboard_name = "production-dashboard"
 # ... configuration
}

Only create in non-production
resource "aws_instance" "debug" {
 count = terraform.workspace != "prod" ? 1 : 0
 # ... debug instance configuration
}
```

## Pattern 4: Workspace-Specific Backend Config

```
hcl

backend.tf
terraform {
 backend "s3" {
 bucket = "my-terraform-state"
 key = "app/terraform.tfstate"
 region = "us-east-1"
 dynamodb_table = "terraform-locks"

 # Workspace prefix creates: env:/dev/, env:/staging/, env:/prod/
 workspace_key_prefix = "env"
 }
}
```

# Best Practices

## 1. Always Verify Current Workspace

bash

```
Before applying changes
tofu workspace show
tofu plan | grep "workspace"
```

## 2. Use Workspace Naming Conventions

bash

```
Good
dev, staging, prod

Also good
dev-useast1, dev-uswest2
feature-auth, feature-payments

Bad
johns-test, temporary, new, old
```

## 3. Document Workspace Strategy

markdown

```
README.md
Workspaces
- `dev`: Development environment (auto-deploys from main)
- `staging`: Pre-production testing (manual deploys)
- `prod`: Production (requires approval)
```

## 4. Implement Safety Checks

bash

```
Add to CI/CD pipeline
if ["$(tofu workspace show)" == "prod"]; then
 echo "Deploying to PRODUCTION. Continue? (yes/no)"
 read confirmation
 ["$confirmation" != "yes"] && exit 1
fi
```

## **Key Limitations**

### **1. No Cross-Workspace References**

- Can't reference resources from other workspaces
- Each workspace is completely isolated

### **2. Workspace State is Mutable**

- Deleting workspace deletes all resource state
- No automatic protection for production workspaces

### **3. Limited Backend Support**

- Not all backends support workspaces
- Some cloud providers have better native solutions

### **4. No Built-in RBAC**

- Everyone with state access can use any workspace
- Need external controls for production protection

## Summary

Workspaces are a powerful feature for managing multiple environments from a single codebase, but they're not a one-size-fits-all solution. Use them when:

- Infrastructure is fundamentally similar across environments
- You need rapid environment creation
- Single team manages all environments
- State isolation is sufficient for your security needs

Consider alternatives when:

- Environments have significantly different architectures
- Different teams need different access controls
- Regulatory requirements demand hard separation
- Environments have different lifecycles

## **Environment Management with Workspaces**

## Overview

This section covers practical patterns for managing dev, staging, and production environments using OpenTofu workspaces. We'll explore configuration strategies, variable management, and real-world deployment patterns.

## Environment Design Philosophy

### The Three-Tier Standard

Most organizations use a three-tier environment structure:

```
Development (dev)
 ↓
Staging (staging/stage)
 ↓
Production (prod)
```

## **Environment Characteristics:**

| Aspect  | Dev                    | Staging             | Production               |
|---------|------------------------|---------------------|--------------------------|
| Purpose | Active development     | Pre-prod testing    | Live users               |
| Size    | Small (cost-optimized) | Medium (prod-like)  | Full scale               |
| Data    | Synthetic/scrambled    | Sanitized prod copy | Real user data           |
| Uptime  | Best effort            | High                | Critical                 |
| Changes | Multiple per day       | Daily/weekly        | Controlled releases      |
| Cost    | Minimize               | Balance             | Optimize for performance |

## Workspace Mapping

hcl

```
Workspace to environment mapping
dev → Development environment
staging → Staging/QA environment
prod → Production environment
```

## **Configuration Strategies**

## Strategy 1: Inline Conditionals

Simple conditional logic directly in resources:

```
hcl

main.tf
resource "aws_instance" "web" {
 ami = var.ami_id
 instance_type = terraform.workspace == "prod" ? "t3.large" : "t3.small"

 monitoring = terraform.workspace == "prod" ? true : false

 tags = {
 Name = "${terraform.workspace}-web"
 Environment = terraform.workspace
 }
}
```

### Pros:

- Simple and direct
- Easy to understand
- No extra files

### Cons:

- Can become complex quickly
- Hard to maintain with many differences
- Limited to simple conditionals

## Strategy 2: Variable Maps

Define environment-specific values in variable maps:

```
hcl

variables.tf
variable "instance_config" {
 type = map(object({
 instance_type = string
 instance_count = number
 volume_size = number
 }))
 default = {
 dev = {
 instance_type = "t3.small"
 instance_count = 1
 volume_size = 20
 }
 staging = {
 instance_type = "t3.medium"
 instance_count = 2
 volume_size = 50
 }
 prod = {
 instance_type = "t3.large"
 instance_count = 5
 volume_size = 100
 }
 }
}
```

```
main.tf
resource "aws_instance" "app" {
 count = var.instance_config[terraform.workspace].instance_count
 ami = var.ami_id
 instance_type = var.instance_config[terraform.workspace].instance_type

 root_block_device {
 volume_size = var.instance_config[terraform.workspace].volume_size
 }

 tags = {
 Name = "${terraform.workspace}-app-${count.index + 1}"
 }
}
```

### Pros:

- Clear separation of configuration
- Easy to add new environments
- Type-safe with object validation

### Cons:

- All values in version control
- Can't override easily per deployment
- Variables file can get large

### Strategy 3: Separate Variable Files

External variable files for each environment:

hcl

```
variables.tf
variable "instance_type" {
 type = string
 description = "EC2 instance type"
}

variable "instance_count" {
 type = number
 description = "Number of instances to create"
}

variable "enable_monitoring" {
 type = bool
 description = "Enable detailed CloudWatch monitoring"
}
```

hcl

```
dev.tfvars
instance_type = "t3.small"
instance_count = 1
enable_monitoring = false
```

hcl

```
staging.tfvars
instance_type = "t3.medium"
instance_count = 2
enable_monitoring = true
```

hcl

```
prod.tfvars
instance_type = "t3.large"
instance_count = 5
enable_monitoring = true
```

bash

```
Usage
tofu workspace select dev
tofu apply -var-file="dev.tfvars"
```

#### Pros:

- Clean separation of concerns
- Easy to override values
- Can use different sources (git, vault, etc.)

#### Cons:

- Need to remember correct var file
- Easy to apply wrong file to wrong workspace
- Extra files to manage

## Strategy 4: Locals with Complex Logic (Recommended)

Combine workspace interpolation with locals for complex configurations:

hcl

```
locals.tf
locals {
 # Environment-specific configurations
 environments = {
 dev = {
 # Compute
 instance_type = "t3.small"
 instance_count = 1
 min_size = 1
 max_size = 2

 # Database
 db_instance_class = "db.t3.micro"
 db_allocated_storage = 20
 db_multi_az = false

 # Networking
 vpc_cidr = "10.0.0.0/16"

 # Features
 enable_monitoring = false
 enable_backup = false
 backup_retention = 1

 # Scaling
 enable_autoscaling = false
 }

 staging = {
 # Compute
 instance_type = "t3.medium"
 instance_count = 2
 min_size = 2
 max_size = 4

 # Database
 db_instance_class = "db.t3.small"
 db_allocated_storage = 50
 db_multi_az = false

 # Networking
 vpc_cidr = "10.1.0.0/16"

 # Features
 enable_monitoring = true
 enable_backup = true
 backup_retention = 7

 # Scaling
 enable_autoscaling = true
 }

 prod = {
 # Compute
 instance_type = "t3.large"
 instance_count = 5
 min_size = 3
 max_size = 10

 # Database
 db_instance_class = "db.r5.large"
 db_allocated_storage = 100
 db_multi_az = true

 # Networking
 vpc_cidr = "10.2.0.0/16"

 # Features
 enable_monitoring = true
 enable_backup = true
 backup_retention = 30

 # Scaling
 enable_autoscaling = true
 }
 }
}
```

## Strategy 4: Locals with Complex Logic (Recommended)

```
Get config for current workspace
env = local.environments[terraform.workspace]

Common tags applied to all resources
common_tags = {
 Environment = terraform.workspace
 ManagedBy = "OpenTofu"
 Project = "webapp"
 CostCenter = terraform.workspace == "prod" ? "productio" : "non-production"
}
}
```

## Strategy 4: Locals with Complex Logic (Recommended)

```
main.tf - Clean resource definitions
resource "aws_instance" "app" {
 count = local.env.instance_count
 ami = var.ami_id
 instance_type = local.env.instance_type

 monitoring = local.env.enable_monitoring

 root_block_device {
 volume_size = 30
 }

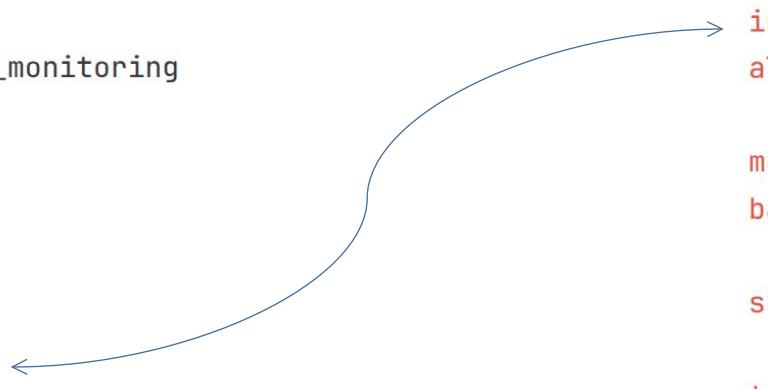
 tags = merge(
 local.common_tags,
 {
 Name = "${terraform.workspace}-app-${count.index + 1}"
 }
)
}
```

```
resource "aws_db_instance" "main" {
 identifier = "${terraform.workspace}-database"
 engine = "postgres"
 engine_version = "15.3"
 instance_class = local.env.db_instance_class
 allocated_storage = local.env.db_allocated_storage

 multi_az = local.env.db_multi_az
 backup_retention_period = local.env.backup_retention

 skip_final_snapshot = terraform.workspace != "prod"

 tags = local.common_tags
}
```



### Strategy 4 Combine workspace

#### Pros:

- Clean separation of config and infrastructure code
- Type-safe and validated
- Easy to see all environment differences
- Supports complex logic
- Single source of truth

#### Cons:

- All values in code (use variables for secrets)
- Requires workspace discipline

# Hands-On Labs: Workspaces & Environment Segregation

## Lab Overview

This hands-on lab series will take you through implementing workspaces for environment segregation.

We'll deploy real AWS infrastructure across dev, staging, and production environments.

### Lab Structure:

**Lab :** Basic workspace usage with simple web servers

### Prerequisites:

- AWS account with admin access
- AWS CLI configured
- OpenTofu  $\geq$  1.6.0 installed
- SSH key pair available
- Basic understanding of AWS (EC2, VPC, S3)

### Scenario

You'll deploy a simple web application across three environments (dev, staging, prod) with different instance counts and sizes per environment.

### Architecture

|              |                                |
|--------------|--------------------------------|
| Development: | 1x t3.small + VPC              |
| Staging:     | 2x t3.medium + VPC             |
| Production:  | 3x t3.large + VPC + Monitoring |

## **OPA Fundamentals - Understanding Policy-as-Code**

## **1. Introduction to Policy-as-Code**

### **What is Policy-as-Code?**

Policy-as-Code is the practice of writing organizational policies, compliance rules, and security standards as executable code rather than documentation. Just as Infrastructure-as-Code (IaC) transformed infrastructure management, Policy-as-Code transforms governance.

# **Why Policy-as-Code Matters**

## **Why Policy-as-Code Matters**

### **Traditional Approach Problems:**

- Manual reviews are slow and error-prone
- Policies documented in wikis get outdated
- Inconsistent enforcement across teams
- No way to test policies before applying them
- Difficult to audit compliance

## **Why Policy-as-Code Matters**

### **Policy-as-Code Benefits:**

- Automated, consistent enforcement
- Policies version controlled alongside infrastructure
- Testable before deployment
- Fast feedback to developers
- Clear audit trail
- Shift-left security (catch issues early)

## Why Policy-as-Code Matters

### Real-World Scenario

Imagine you're managing infrastructure for a healthcare company:

- HIPAA requires encryption at rest and in transit
- Finance requires cost controls on resources
- Security requires specific tagging for compliance tracking
- Architecture requires resources in approved regions only

Without Policy-as-Code:

- Manual review of every pull request (slow)
- Mistakes slip through (risky)
- Policies drift between teams (inconsistent)
- Violations discovered after deployment (expensive)

## **Why Policy-as-Code Matters**

With Policy-as-Code:

- Automatic validation on every commit
- Instant feedback to developers
- Consistent enforcement across all teams
- Issues caught before any resources are created

## 2. Open Policy Agent (OPA) Architecture

### What is OPA?

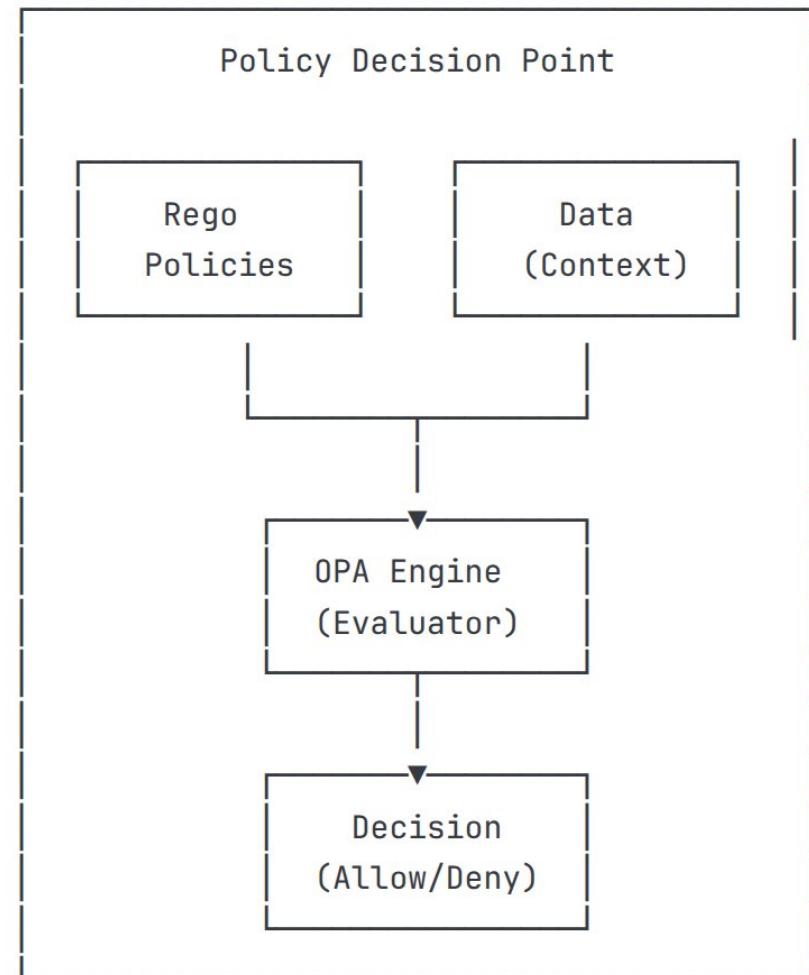
Open Policy Agent is a general-purpose policy engine that:

- Evaluates policies written in the Rego language
- Works with any JSON/YAML data
- Provides a unified way to enforce policies
- Can be embedded in applications or used as a service

### Key Components:

1. **Rego Policies:** Rules written in the Rego language
2. **Data/Context:** Information about what's being evaluated
3. **OPA Engine:** Evaluates policies against data
4. **Decision:** Result (allow/deny) with reasoning

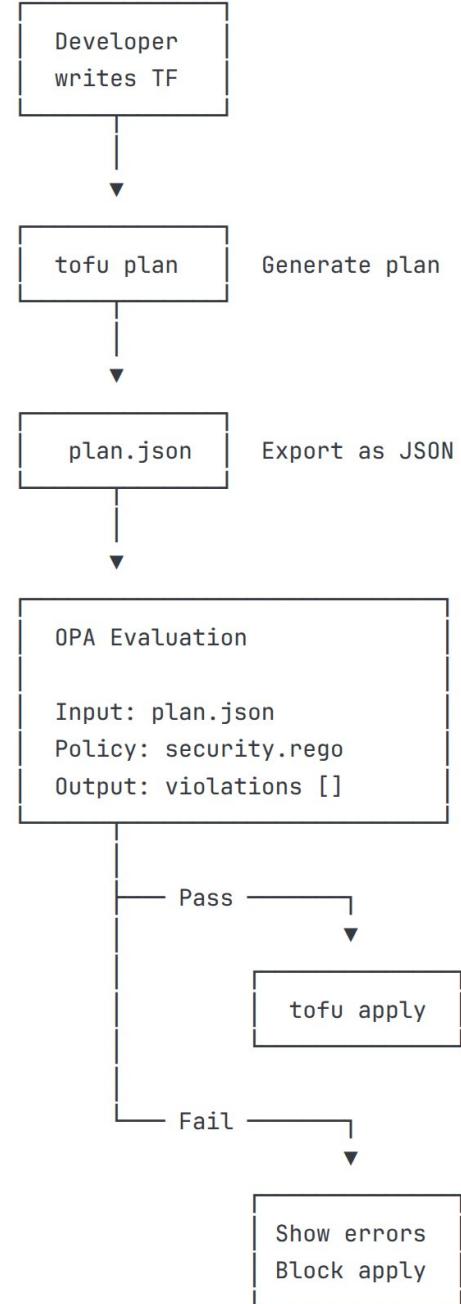
### OPA Components



## How OPA Works with OpenTofu

### Policy Evaluation Workflow

1. **Generate Plan:** OpenTofu creates execution plan
2. **Export JSON:** Convert plan to JSON format
3. **Load Policy:** OPA loads relevant Rego policies
4. **Evaluate:** OPA checks plan against policies
5. **Report:** OPA returns violations or approval
6. **Decide:** Allow apply or require fixes



### 3. OPA vs Other Policy Tools

Comparison Matrix

| Feature        | OPA                           | Sentinel       | Cloud Custodian |
|----------------|-------------------------------|----------------|-----------------|
| Language       | Rego                          | HCL-like       | YAML            |
| Cloud-agnostic | Yes                           | No (HashiCorp) | Partial         |
| Runtime        | Wandora, Terraform Enterprise | Enterprise     | Cloud-specific  |
| Testing        | Built-in                      | Limited        | Limited         |
| Complexity     | Moderate                      | Low            | Low-Moderate    |
| Flexibility    | Very High                     | Moderate       | Moderate        |

## How OPA Works with OpenTofu

## When to Use OPA

### **Best For:**

- Multi-cloud environments
- Complex policy logic
- Microservices authorization
- Kubernetes admission control
- Organizations wanting open-source solutions

### **Consider Alternatives If:**

- Using only HashiCorp Terraform Enterprise (Sentinel)
- Simple policies only (might be overkill)
- Team unfamiliar with logic programming

## 4. Understanding the Rego Language

### What is Rego?

Rego (pronounced "ray-go") is a declarative language designed for policy writing:

- Based on Datalog (logic programming)
- Focused on querying and transforming data
- Optimized for policy evaluation
- Similar to SQL in purpose, different in syntax

### How OPA Works with OpenTofu

## Rego Philosophy

### Declarative, Not Imperative:

```
rego

Rego: Declare what should be true
deny[msg] {
 resource := input.resource_changes[_]
 resource.type == "aws_s3_bucket"
 not resource.change.after.encryption
 msg := "S3 buckets must be encrypted"
}
```

```
Versus imperative (pseudocode):
for each resource in plan:
if resource.type == "s3_bucket":
if not resource.encrypted:
return "error"
```

### Key Characteristics:

- Rules evaluate to true or false
- No variables in traditional sense (only bindings)
- No side effects or state changes
- Order-independent execution

## How OPA Works with OpenTofu

## 5. Setting Up OPA with OpenTofu

### Installation

Install OPA:

# Linux

```
curl -L -o opa https://openpolicyagent.org/downloads/latest/opa_linux_amd64
chmod +x opa
sudo mv opa /usr/local/bin/
```

# macOS

```
brew install opa
```

# Windows

```
curl -L -o opa.exe https://openpolicyagent.org/downloads/latest/opa_windows_amd64.exe
```

# Verify

```
opa version
```

## Directory Structure

```
project/
 └── infrastructure/
 ├── main.tf
 ├── variables.tf
 └── outputs.tf
 └── policies/
 ├── tagging.rego
 ├── security.rego
 ├── cost.rego
 └── test/
 ├── tagging_test.rego
 ├── security_test.rego
 └── cost_test.rego
 └── scripts/
 ├── check-policies.sh
 └── generate-plan.sh
```

## Basic Workflow Script

bash

```
#!/bin/bash
scripts/check-policies.sh

set -e

echo "⌚ Generating OpenTofu plan..."
cd infrastructure
tofu plan -out=tfplan.binary

echo "📄 Converting plan to JSON..."
tofu show -json tfplan.binary > tfplan.json

echo "🔍 Evaluating policies..."
cd ../policies
opa eval --data . --input ../infrastructure/tfplan.json \
--format pretty \
"data.terraform.deny"

if [$? -eq 0]; then
 echo "✅ All policies passed!"
 exit 0
else
 echo "❗ Policy violations found!"
 exit 1
fi
```

## 6. OpenTofu Plan Structure

Understanding the JSON plan structure is crucial for writing policies.

### Plan JSON Overview

```
json
{
 "format_version": "1.2",
 "terraform_version": "1.6.0",
 "planned_values": {
 "root_module": {
 "resources": [...]
 }
 },
}
```

### Key Sections for Policy Writing

**resource\_changes:** Most important for policies - contains all resource modifications:

- **actions**: ["create"], ["update"], ["delete"], ["create", "delete"] (replace)
- **before**: State before changes (null for new resources)
- **after**: State after changes (null for deleted resources)
- **type**: Resource type (aws\_instance, aws\_s3\_bucket, etc.)

```
"resource_changes": [
 {
 "address": "aws_instance.web",
 "mode": "managed",
 "type": "aws_instance",
 "name": "web",
 "provider_name": "registry.opentofu.org/hashicorp/aws",
 "change": {
 "actions": ["create"],
 "before": null,
 "after": {
 "ami": "ami-12345678",
 "instance_type": "t3.micro",
 "tags": {
 "Name": "web-server"
 }
 },
 "after_unknown": {...}
 }
],
 "configuration": {...}
}
```

## Example Access in Rego:

---

```
rego

Get all resources being created
resource := input.resource_changes[_]
resource.change.actions[_] == "create"

Get resource type
resource.type == "aws_s3_bucket"

Get configuration values
resource.change.after.instance_type == "t3.large"
```

```
package aws.compliance

import rego.v1

Helper: Define allowed types set
allowed_types := {"t2.micro", "t3.micro"}

Rule 1: Restrict EC2 Instance Types
deny contains msg if {
 # Find all resources in the plan
 resource := input.resource_changes[_]

 # Filter for EC2 instances only
 resource.type == "aws_instance"

 # Check if the instance type is in the allowed list
 not allowed_types[resource.change.after.instance_type]

 # Return a specific error message
 msg := sprintf(
 "Invalid instance type: '%v'. Allowed types are: %v",
 [resource.change.after.instance_type, allowed_types]
)
}

Rule 2: Enforce Mandatory 'Environment' Tag
deny contains msg if {
 resource := input.resource_changes[_]
 resource.type == "aws_instance"

 # Check if 'Environment' tag exists
 not resource.change.after.tags.Environment

 msg := "EC2 instance is missing the mandatory 'Environment' tag."
}
```

### 3. How it Works (Evaluation)

If you evaluate the **Policy** against the **Input**, OPA will return two violations:

1. **Violation 1:** The input uses `t2.large`, but the policy only allows `t2.micro` or `t3.micro`.
2. **Violation 2:** The input has a `Name` tag, but is missing the required `Environment` tag.

### How to test this locally

If you have the OPA binary installed, you can run this command in your terminal:

Bash

```
opa eval --data policy.rego --input input.json "data.aws.compliance.deny"
```

```
● jilg@node-214:~/tofu_infra/lesson17_OPA$ opa eval --data policy.rego --input input.json "data.aws.compliance.deny"
{
 "result": [
 {
 "expressions": [
 {
 "value": [
 "EC2 instance is missing the mandatory 'Environment' tag.",
 "Invalid instance type: 't2.large'. Allowed types are: {\\"t2.micro\\", \\"t3.micro\\\"}"
],
 "text": "data.aws.compliance.deny",
 "location": {
 "row": 1,
 "col": 1
 }
 }
]
 }
]
}
```

**Rego Language Basics - Writing Effective Policies**

**Refer - Rego\_language\_tutorial.html**

# Hands-On Lab: OPA Policies for Security Compliance

## Lab Overview

**Difficulty:** Intermediate

### Objectives:

Write OPA policies for real-world compliance requirements  
Test policies against OpenTofu plans  
Fix policy violations in infrastructure code  
Integrate policies into development workflow

### 2. Verify Installation:

---

```
bash
tofu version # Should show v1.6.0 or higher
opa version # Should show v0.60.0 or higher
aws --version # Should show aws-cli/2.x
jq --version # Should show jq-1.6 or higher
```

## Lab Environment Setup

### Prerequisites

#### Tools Installed:

OpenTofu CLI  
OPA CLI  
AWS CLI (configured)  
jq  
Git

lesson19-OPA\_Policies\_lab

## **OpenTofu + Spacelift Integration: Deploy an EC2 Instance**

# Integrating OpenTofu with Spacelift and deploying an AWS EC2 instance.

## Overview

### What you'll accomplish:

- Set up a Spacelift account and connect it to your VCS (GitHub/GitLab)
- Configure AWS credentials using Spacelift's native cloud integration
- Write OpenTofu code to provision an EC2 instance
- Create a Spacelift Stack and deploy infrastructure

### Architecture Flow:

GitHub Repo → Spacelift Stack → AWS Cloud Integration → EC2 Instance

### Prerequisites

1. **Spacelift Account** - Sign up at [spacelift.io](https://spacelift.io) (14-day free trial available)
2. **AWS Account** - With permissions to create IAM roles and EC2 instances
3. **GitHub/GitLab Account** - To host your OpenTofu code
4. **Git** - To push your code to the repository

# **GraphQL Fundamentals**

## Understanding GraphQL Queries

GraphQL is a query language for APIs that allows you to request exactly the data you need:

### Key Concepts:

- **Queries:** Read operations to fetch data
- **Mutations:** Write operations to modify data
- **Fields:** Specific data points you want to retrieve
- **Arguments:** Parameters to filter or customize results

### Basic Structure:

```
graphql
query {
 stacks {
 id
 name
 state
 }
}
```

## Spacelift GraphQL Schema

Common objects you'll work with:

1. **Stack:** Infrastructure stack configuration

- Fields: `id`, `name`, `state`, `repository`, `branch`

2. **Run:** Deployment execution

- Fields: `id`, `state`, `type`, `createdAt`, `triggeredBy`

3. **Viewer:** Current authenticated user

- Fields: `username`, `permissions`

## **Migration Planning - Detailed Concepts**

1. Why Migrate to OpenTofu?
2. Compatibility Matrix
3. Dependency Identification
4. Migration Checklist Development
5. Implementation Steps

# Why Migrate to OpenTofu?

## The Licensing Shift

### What Happened?

In August 2023, HashiCorp changed Terraform's license from the Mozilla Public License v2.0 (MPL 2.0) to the Business Source License v1.1 (BSL 1.1). This change had significant implications for the infrastructure-as-code community.

### Key License Differences:

| Aspect               | MPL 2.0 (Old)    | BSL 1.1 (New)                     |
|----------------------|------------------|-----------------------------------|
| Type                 | True open source | Source-available                  |
| Commercial Use       | Unrestricted     | Restricted for competitors        |
| Derivative Works     | Allowed freely   | Prohibited for competing products |
| Community Governance | Open             | Corporate controlled              |
| Future Certainty     | Predictable      | Subject to change                 |

### **What BSL 1.1 Means for Users:**

The Business Source License prohibits using Terraform to create competing offerings. While most infrastructure teams are safe, the license creates uncertainty:

- **Direct Impact:** Companies building Terraform-as-a-Service platforms
- **Indirect Impact:** Organizations concerned about future restrictions
- **Strategic Risk:** Dependency on a single vendor's licensing decisions
- **Vendor Lock-in:** Reduced leverage in negotiations with HashiCorp

## **Open-Source Continuity**

### **The OpenTofu Initiative**

OpenTofu emerged as a community response to Terraform's licensing change, committed to:

1. **True Open Source:** MPL 2.0 license permanently
2. **Community Governance:** Linux Foundation management
3. **Backward Compatibility:** Drop-in replacement for Terraform
4. **Transparent Development:** Public roadmap and decision-making
5. **Vendor Neutrality:** No single company control

## **Who's Behind OpenTofu?**

- **Linux Foundation:** Neutral governance body
- **Community Contributors:** Thousands of developers
- **Enterprise Sponsors:** Spacelift, env0, Scalr, Gruntwork, and others
- **Cloud Providers:** Support from major cloud platforms

## **Long-term Sustainability:**

OpenTofu's governance model ensures:

- No surprise licensing changes
- Community-driven feature development
- Multiple commercial support options
- Reduced vendor dependency risk

## **Business Risk Mitigation**

### **Strategic Considerations for Migration:**

#### **1. Licensing Risk Reduction**

- Remove dependency on proprietary licensing terms
- Eliminate concerns about future license changes
- Maintain full commercial use rights
- Avoid potential compliance issues

#### **2. Vendor Independence**

- Multiple commercial support options available
- No single point of control
- Competitive marketplace for tools and services
- Reduced negotiation leverage concerns

### **3. Community Investment**

- Contributions benefit entire community
- Transparent feature development
- Shared maintenance burden
- Collaborative problem-solving

### **4. Future-Proofing**

- Guaranteed open-source access
- Community continuity regardless of corporate changes
- Multiple implementation options
- Reduced technical debt risk

## **5. Cost Considerations**

- No licensing fees for OpenTofu
- Competitive ecosystem for commercial support
- Open marketplace for tools and integrations
- Reduced vendor lock-in costs

## **Feature Parity and Development**

### **Current State (OpenTofu 1.6+):**

OpenTofu maintains feature parity with Terraform v1.5.x and adds:

- All core Terraform functionality
- Improved error messages and debugging
- Enhanced state encryption options
- Community-requested features
- Performance optimizations

## **Future Development:**

OpenTofu's roadmap includes:

- Enhanced module testing capabilities
- Improved provider development experience
- Additional state backend options
- Performance improvements
- Community-driven feature additions

## **Migration Timing Considerations:**

### **Migrate Now If:**

- Concerned about licensing uncertainty
- Value open-source governance
- Want to support community development
- Planning new infrastructure projects
- Updating existing tooling anyway

### **Wait If:**

- Using Terraform Cloud features heavily (consider alternatives)
- Dependent on brand-new Terraform features (check compatibility)
- Organization has strict change control (plan carefully)
- Terraform Enterprise with specific version requirements

## Compatibility Matrix

### Version Compatibility

#### Critical Compatibility Facts:

OpenTofu 1.6.0+ is designed as a drop-in replacement for Terraform v1.5.x and earlier.

Understanding the compatibility matrix is essential for migration planning.

#### Compatibility Table:

| Component              | Terraform          | OpenTofu | Compatibility Status     |
|------------------------|--------------------|----------|--------------------------|
| <b>HCL Syntax</b>      | All versions       | v1.6.0+  | 100% compatible          |
| <b>State Format</b>    | v1.5.x and earlier | v1.6.0+  | Fully compatible         |
| <b>Providers</b>       | All versions       | v1.6.0+  | Works with all providers |
| <b>Modules</b>         | All versions       | v1.6.0+  | No changes needed        |
| <b>Backend Configs</b> | All versions       | v1.6.0+  | Same syntax              |
| <b>CLI Commands</b>    | v1.x               | v1.x     | Identical interface      |

## HCL (HashiCorp Configuration Language) Compatibility

### 100% Backward Compatible

All valid Terraform HCL configuration files work identically with OpenTofu. No syntax changes required.

```
This configuration works identically in Terraform and OpenTofu

resource "aws_instance" "web" {
 ami = data.aws_ami.ubuntu.id
 instance_type = var.instance_type

 tags = {
 Name = "web-server"
 Environment = var.environment
 ManagedBy = "OpenTofu" # Can update this, but not required
 }
}

variable "instance_type" {
 description = "EC2 instance type"
 type = string
 default = "t2.micro"

 validation {
 condition = contains(["t2.micro", "t2.small", "t3.micro"], var.instance_type)
 error_message = "Instance type must be t2.micro, t2.small, or t3.micro."
 }
}
```

## **Advanced Features Compatibility:**

All Terraform HCL features work in OpenTofu:

-  Dynamic blocks
-  For expressions
-  Conditional expressions
-  Function calls
-  Type constraints
-  Validation rules
-  Sensitive values
-  Lifecycle rules
-  Provisioners
-  Local values

## Provider Compatibility

### Provider Sources Continue to Work

OpenTofu maintains compatibility with all Terraform providers, including those from the HashiCorp registry.

#### Provider Source Resolution:

hcl

```
terraform {
 required_providers {
 # Both HashiCorp and OpenTofu registries work
 aws = {
 source = "hashicorp/aws" # ✓ Works perfectly
 version = "~> 5.0"
 }

 random = {
 source = "hashicorp/random" # ✓ Also works
 version = "~> 3.5"
 }

 # OpenTofu registry (optional, identical providers)
 custom = {
 source = "opentofu/custom" # ✓ OpenTofu registry
 version = "~> 1.0"
 }
 }
}
```

## **Provider Registry Behavior:**

1. **Default Registry:** OpenTofu uses HashiCorp's provider registry by default
2. **No Changes Needed:** Existing `source` declarations work as-is
3. **Mirror Support:** Can configure provider mirrors if desired
4. **Version Locking:** `.terraform.lock.hcl` works identically

## **Important Notes:**

- No need to change provider sources during migration
- Provider plugins are identical (same binaries work)
- Provider version constraints work the same way
- Provider authentication methods unchanged

## State File Compatibility

### State Format is Identical

The `.tfstate` file format is 100% compatible between Terraform and OpenTofu.

```
{
 "version": 4,
 "terraform_version": "1.5.7", // This field may show different versions
 "serial": 1,
 "lineage": "a3c5e8f1-8c7a-4d9b-b2e6-1f5d7c9a3b4e",
 "outputs": {},
 "resources": [
 {
 "mode": "managed",
 "type": "aws_instance",
 "name": "example",
 "provider": "provider[\"registry.terraform.io/hashicorp/aws\"]",
 "instances": [
 {
 "schema_version": 1,
 "attributes": {
 "ami": "ami-0c55b159cbfafe1f0",
 "instance_type": "t2.micro"
 }
 }
]
 }
]
}
```

## **State Migration Process:**

When you run `tofu init -migrate-state`:

1. **Reads Existing State:** OpenTofu reads the Terraform state file
2. **Validates Format:** Confirms state format compatibility
3. **Preserves Data:** Maintains all resource tracking information
4. **Updates Metadata:** Updates version field to reflect OpenTofu
5. **Maintains Lineage:** Preserves state lineage for tracking

## State Backend Compatibility:

All Terraform backends work with OpenTofu:

| Backend Type           | Compatibility | Notes                                    |
|------------------------|---------------|------------------------------------------|
| <b>Local</b>           | Perfect       | Default, no changes needed               |
| <b>S3</b>              | Perfect       | Same bucket configuration                |
| <b>Azure Storage</b>   | Perfect       | Identical settings                       |
| <b>GCS</b>             | Perfect       | Same bucket syntax                       |
| <b>Consul</b>          | Perfect       | No modifications required                |
| <b>Kubernetes</b>      | Perfect       | Same secret configuration                |
| <b>HTTP</b>            | Perfect       | REST API unchanged                       |
| <b>Terraform Cloud</b> | Limited       | Use alternative (Spacelift, env0, Scalr) |

## Module Compatibility

### Modules Work Without Modification

All Terraform modules work identically with OpenTofu.

```
module "vpc" {
 # Local modules
 source = "./modules/vpc" # Works perfectly
}

module "security" {
 # Terraform Registry
 source = "terraform-aws-modules/security-group/aws"
 version = "5.1.0" # Works perfectly
}

module "network" {
 # Git sources
 source = "git::https://github.com/org/terraform-modules.git//networking?ref=v1.2.0" # Works
}

module "storage" {
 # GitHub shorthand
 source = "github.com/org/modules//storage" # Works
}

module "compute" {
 # S3 bucket
 source = "s3::https://s3-eu-west-1.amazonaws.com/bucket/modules/compute.zip" # Works
}
```

## **Publishing Modules:**

You can publish modules to:

- **Terraform Registry:** Still works with OpenTofu consumers
- **OpenTofu Registry:** Designed for OpenTofu (same format)
- **Private Registries:** Organization-specific module hosting
- **Git Repositories:** Direct source references

## CLI Command Compatibility

### Identical Command-Line Interface

OpenTofu uses `tofu` instead of `terraform`, but all commands and flags work identically.

#### Command Mapping:

| Terraform Command               | OpenTofu Command           | Notes                |
|---------------------------------|----------------------------|----------------------|
| <code>terraform init</code>     | <code>tofu init</code>     | Identical behavior   |
| <code>terraform plan</code>     | <code>tofu plan</code>     | Same options         |
| <code>terraform apply</code>    | <code>tofu apply</code>    | Same workflow        |
| <code>terraform destroy</code>  | <code>tofu destroy</code>  | Identical            |
| <code>terraform output</code>   | <code>tofu output</code>   | Same format          |
| <code>terraform state</code>    | <code>tofu state</code>    | All subcommands work |
| <code>terraform import</code>   | <code>tofu import</code>   | Identical syntax     |
| <code>terraform validate</code> | <code>tofu validate</code> | Same validation      |
| <code>terraform fmt</code>      | <code>tofu fmt</code>      | Identical formatting |

## Migration Simplicity Example:

---

bash

*# Before migration (Terraform)*

```
terraform init
terraform plan
terraform apply
```

*# After migration (OpenTofu) - literally just change the binary name*

```
tofu init
tofu plan
tofu apply
```

## **Dependency Identification**

### **Identifying Module Dependencies**

### **Why Module Identification Matters**

Understanding your module dependencies is critical for:

- Ensuring continued functionality after migration
- Planning module registry updates
- Validating version constraints
- Documenting infrastructure architecture

## Discovering Module Dependencies:

### 1. Configuration File Analysis

```
bash
```

```
Find all module blocks in your configuration
grep -r "^module " . --include="*.tf"

Find module sources
grep -r "source\s*=" . --include="*.tf" | grep "module"
```

## Discovering Module Dependencies:

### 2. Review Each Module Declaration

---

hcl

```
module "vpc" {
 source = "terraform-aws-modules/vpc/aws" # Public registry module
 version = "5.1.2" # Version constraint

 # Note module configuration for validation testing
 name = "my-vpc"
 cidr = "10.0.0.0/16"
}

module "custom_networking" {
 source = "./modules/networking" # Local module (ensure code is present.

 vpc_id = module.vpc.vpc_id
}

module "private_module" {
 source = "git::https://github.com/myorg/private-modules.git//networking?ref=v2.1.0"
 # Git module (ensure access credentials work)
}
```

### **3. Document Module Inventory**

Create a module dependency matrix:

| Module Name | Source Type | Source Location                  | Version | Used By       | Critical? |
|-------------|-------------|----------------------------------|---------|---------------|-----------|
| vpc         | Registry    | terraform-aws-modules/vpc/aws    | 5.1.2   | main.tf       | Yes       |
| networking  | Local       | ./modules/networking             | N/A     | main.tf       | Yes       |
| private_net | Git         | github.com/org/modules           | v2.1.0  | network.tf    | Yes       |
| monitoring  | Registry    | terraform-aws-modules/cloudwatch | 4.0.1   | monitoring.tf | No        |

## Provider Dependency Mapping

### Provider Version Constraints

Identify all providers and their version constraints:

#### 1. Locate Provider Declarations

---

bash

```
Find all required_providers blocks
grep -A 10 "required_providers" . -r --include="*.tf"
```

## 2. Document Provider Requirements

---

hcl

```
terraform {
 required_version = ">= 1.5.0" # Minimum Terraform version

 required_providers {
 aws = {
 source = "hashicorp/aws"
 version = "~> 5.0" # Document this constraint
 }

 random = {
 source = "hashicorp/random"
 version = ">= 3.5.0" # And this one
 }

 tls = {
 source = "hashicorp/tls"
 version = "~> 4.0" # And this one
 }
 }
}
```

### **3. Provider Inventory Table**

| Provider | Source           | Version Constraint | Purpose         | Configuration Needs |
|----------|------------------|--------------------|-----------------|---------------------|
| aws      | hashicorp/aws    | $\sim > 5.0$       | Primary cloud   | Region, credentials |
| random   | hashicorp/random | $\geq 3.5.0$       | Random values   | None                |
| tls      | hashicorp/tls    | $\sim > 4.0$       | SSH keys, certs | None                |

### **4. Validate Provider Availability**

All HashiCorp providers work with OpenTofu. Check that:

- Provider versions are available in the registry
- Version constraints are not too restrictive
- Provider configuration will work post-migration

## Remote State Backend Dependencies

### State Backend Configuration Review

#### 1. Identify Current Backend

```
bash
```

```
Check for backend configuration
grep -A 10 "backend \"\" . -r --include="*.tf"
```

#### 2. Document Backend Settings

```
hcl
```

```
terraform {
 backend "s3" {
 bucket = "my-terraform-state" # S3 bucket name
 key = "prod/vpc/terraform.tfstate" # State file path
 region = "us-east-1" # AWS region
 encrypt = true # Encryption enabled
 dynamodb_table = "terraform-locks" # Lock table

 # Note these for post-migration validation
 }
}
```

### **3. Backend Dependency Checklist**

- Backend type identified (S3, Azure Storage, GCS, etc.)
- Access credentials documented
- Encryption settings noted
- Lock mechanism identified
- Backup procedure confirmed
- Alternative backends evaluated (if using Terraform Cloud)

## 4. Special Considerations for Terraform Cloud

If using Terraform Cloud backend:

```
hcl

terraform {
 backend "remote" {
 organization = "my-org"

 workspaces {
 name = "production"
 }
 }
}
```

### Migration Options:

- **Option 1:** Migrate to equivalent service (Spacelift, env0, Scalr)
- **Option 2:** Switch to standard remote backend (S3, GCS, Azure)
- **Option 3:** Use local backend temporarily for testing

## Third-Party Tool Integrations

### Identify Ecosystem Dependencies

#### 1. CI/CD Pipeline Tools

Document tools that execute Terraform commands:

| Tool           | Integration Type                | Changes Needed                          |
|----------------|---------------------------------|-----------------------------------------|
| GitHub Actions | Workflow uses terraform CLI     | Update binary name to <code>tofu</code> |
| GitLab CI      | Pipeline script calls terraform | Update commands                         |
| Jenkins        | Terraform plugin or CLI         | Update plugin or scripts                |
| CircleCI       | Docker image with terraform     | Use OpenTofu image                      |
| Azure DevOps   | Terraform task                  | Replace with script task                |

#### Example GitHub Actions Change:

yaml

*# Before*

```
- name: Terraform Plan
 run: terraform plan
```

*# After*

```
- name: OpenTofu Plan
 run: tofu plan
```

## 2. Infrastructure Platforms

| Platform        | Integration             | Migration Path                     |
|-----------------|-------------------------|------------------------------------|
| Spacelift       | Native OpenTofu support | Enable OpenTofu in stack settings  |
| env0            | Native OpenTofu support | Switch tool version in environment |
| Atlantis        | OpenTofu compatible     | Update configuration               |
| Terraform Cloud | Proprietary             | Migrate to alternative             |

### 3. Development Tools

| Tool                        | Purpose                         | Migration Impact           |
|-----------------------------|---------------------------------|----------------------------|
| VS Code Terraform Extension | Syntax highlighting, completion | Works with OpenTofu        |
| IntelliJ Terraform Plugin   | IDE integration                 | Works with OpenTofu        |
| TFLint                      | Linter                          | Works with OpenTofu        |
| Terraform-docs              | Documentation generator         | Works with OpenTofu        |
| Checkov                     | Security scanning               | Works with OpenTofu        |
| Infracost                   | Cost estimation                 | OpenTofu support available |

### 4. Monitoring and Observability

| Tool                          | Purpose            | Compatibility       |
|-------------------------------|--------------------|---------------------|
| Datadog Terraform Integration | Monitoring         | Works with OpenTofu |
| Terraform Compliance          | Policy as code     | Compatible          |
| OPA/Sentinel                  | Policy enforcement | Works with OpenTofu |

# **Migration Checklist Development**

## 1. Environment Assessment

### Infrastructure Inventory:

- List all Terraform projects/workspaces
- Identify Terraform version for each project
- Document provider versions
- Map module dependencies
- List remote state backends
- Identify shared modules

### Example Inventory Spreadsheet:

| Project Name | Terraform Version | Providers       | Modules         | Backend | Priority |
|--------------|-------------------|-----------------|-----------------|---------|----------|
| prod-vpc     | 1.5.7             | aws, random     | vpc, networking | S3      | Critical |
| dev-compute  | 1.5.5             | aws, tls        | ec2, monitoring | S3      | Medium   |
| staging-db   | 1.5.7             | aws, postgresql | rds, backup     | S3      | High     |

## 2. Compatibility Verification

### Version Compatibility Check:

- Confirm Terraform version is 1.5.x or earlier
- Verify provider versions are current
- Check module versions for compatibility
- Review HCL syntax for deprecated features
- Validate state file format version

### Validation Commands:

bash

```
Check Terraform version
terraform version
```

```
Verify providers
terraform providers
```

```
Validate configuration
terraform validate
```

```
Check state file version
grep "version" terraform.tfstate | head -1
```

### 3. Backup Strategy

#### **State File Backup:**

- Create backup of `.tfstate` file
- Backup `.tfstate.backup` if exists
- Document state file location (local or remote)
- Verify backup integrity
- Store backups in version control or secure storage
- Document backup restoration procedure

**Backup Commands:**

*# Local state backup*

```
cp terraform.tfstate terraform.tfstate.pre-migration-$(date +%Y%m%d)
cp terraform.tfstate.backup terraform.tfstate.backup.pre-migration-$(date +%Y%m%d)
```

*# Remote state backup (S3 example)*

```
aws s3 cp s3://my-bucket/terraform.tfstate ./backups/terraform.tfstate.pre-migration-$(date
+%Y%m%d)
```

**Backup Commands:**

*# Create backup directory with all files*

```
tar -czf terraform-backup-$(date +%Y%m%d).tar.gz .terraform terraform.tfstate* *.tf *.tfvars
```

## **Configuration Backup:**

- Commit all configuration files to Git
- Tag repository with pre-migration version
- Create feature branch for migration
- Document Git commit hash
- Archive configuration files separately

**# Git workflow**

```
git add .
git commit -m "Pre-migration snapshot"
git tag -a pre-migration-$(date +%Y%m%d) -m "Snapshot before OpenTofu migration"
git push origin --tags
```

**# Create migration branch**

```
git checkout -b opentofu-migration
```

## 4. Documentation Requirements

### Document Current State:

- Current infrastructure diagram
- Resource inventory
- Dependencies map
- Provider configuration
- Backend configuration
- Workspace structure
- Environment variables

```
Migration Plan: [Project Name]

Current State
- Terraform Version: 1.5.7
- Infrastructure: VPC, EC2, RDS
- State Backend: S3 (bucket: my-terraform-state)
- Last Applied: 2024-01-15
```

```
Migration Timeline
- Planning: Day 1
- Testing: Day 2-3
- Execution: Day 4
- Validation: Day 5
```

```
Rollback Criteria
- State file corruption
- Resource drift detected
- Provider compatibility issues
- Critical errors in plan
```

```
Success Criteria
- `tofu plan` shows no changes
- All resources in state match actual
- Outputs match expected values
- Integration tests pass
```

## 5. Testing Environment Setup

### **Create Safe Testing Environment:**

- Clone production infrastructure to dev/test
- Create isolated state file
- Test migration in non-production first
- Document test results
- Validate rollback procedures

### **Test Environment Checklist:**

- Same Terraform version as production
- Equivalent infrastructure complexity
- Separate AWS account or region
- Independent state backend
- Test credentials configured

## **Migration Execution Phase**

## 1. OpenTofu Installation

### Installation Steps:

- Download OpenTofu binary or use package manager
- Verify binary checksum
- Install OpenTofu (system or user level)
- Verify installation with `tofu version`
- Update PATH if necessary
- Test OpenTofu with simple configuration

```
Verify OpenTofu installation
tofu version
Expected output: OpenTofu v1.6.0 (or higher)

Verify Terraform still available (for rollback)
terraform version
Expected output: Terraform v1.5.x
```

## 2. State Migration

### Migration Steps:

- Navigate to project directory
- Remove Terraform lock file (will be regenerated)
- Run `tofu init -migrate-state`
- Answer 'yes' to migration prompts
- Verify new lock file created
- Confirm state migration in output

```
Remove old lock file (will be regenerated)
rm .terraform.lock.hcl

Initialize with OpenTofu and migrate state
tofu init -migrate-state
Expected prompt: "Do you want to migrate state from Terraform to OpenTofu?"
Answer: yes

Verify lock file regenerated
ls -la .terraform.lock.hcl
```

### 3. Plan Validation

#### Validation Steps:

- Run `tofu plan` to generate new plan
- Review plan output thoroughly
- Verify plan shows zero changes
- Check for any warnings or notices
- Save plan output for comparison
- Compare with pre-migration plan

```
Generate plan with OpenTofu
tofu plan -out=post-migration.plan

Expected output should show:
"No changes. Your infrastructure matches the configuration."

If changes are shown, investigate before proceeding:
- Resource replacements indicate issues
- Attribute changes may indicate provider differences
- Output changes are generally harmless
```

## 4. Resource Verification

### Verification Checklist:

- All resources present in state
- Resource counts match expectations
- No unexpected additions or removals
- Dependencies correctly mapped
- Outputs generate correctly

```
List all resources
tofu state list

Show specific resource details
tofu state show aws_vpc.main

Verify outputs
tofu output

Check resource count
tofu state list | wc -l
```

## Post-Migration Phase

### 1. State Consistency Validation

#### Validation Steps:

- Run `tofu plan` again
- Confirm zero changes shown
- Validate remote state if applicable
- Check state lock functionality
- Verify state backup created
- Test state refresh

```
Refresh state from actual infrastructure
tofu refresh

Plan should still show no changes
tofu plan

For remote state, verify backend access
tofu state pull > current-state.json
cat current-state.json | jq '.serial' # Should match expected version
```

## 2. Module Registry Validation

### Registry Check:

- Verify all modules still accessible
- Check module versions match constraints
- Test module initialization
- Validate private registry access if applicable
- Confirm module outputs still work

### Module Validation:

```
Re-initialize to verify module access
```

```
rm -rf .terraform/modules
```

```
tofu init
```

```
Should download all modules successfully
```

```
Check module versions
```

```
cat .terraform/modules/modules.json | jq '.Modules[] | {key: .Key, source: .Source, version: .Version}'
```

### 3. Provider Functionality Testing

#### Provider Tests:

- Verify provider authentication works
- Test resource creation (in test environment)
- Test resource updates
- Test resource deletion
- Validate provider-specific features
- Check provider version constraints

```
In test environment, create a simple resource
cat > test.tf << 'EOF'
resource "random_pet" "test" {
 length = 2
}

output "pet_name" {
 value = random_pet.test.id
}
EOF

tofu init
tofu apply -auto-approve
tofu destroy -auto-approve
```

## 4. Reference and Documentation Updates

### Update Documentation:

- Update README files (terraform → tofu)
- Modify runbooks and procedures
- Update CI/CD documentation
- Change team wiki/knowledge base
- Update onboarding materials
- Revise troubleshooting guides

# Infrastructure Management

<!-- Before -->

## Prerequisites

- Terraform v1.5+
- AWS CLI configured

## Usage

terraform init

terraform plan

terraform apply

<!-- After -->

## Prerequisites

- OpenTofu v1.6+
- AWS CLI configured

## Usage

tofu init

tofu plan

tofu apply

### Update CI/CD References:

- GitHub Actions workflows
- GitLab CI pipelines
- Jenkins jobs
- CircleCI configurations
- Azure DevOps pipelines
- Custom scripts and tools

```
.github/workflows/infrastructure.yml
```

```
Before
- name: Setup Terraform
 uses: hashicorp/setup-terraform@v2
 with:
 terraform_version: 1.5.7

- name: Terraform Init
 run: terraform init

After
- name: Setup OpenTofu
 uses: opentofu/setup-opentofu@v1
 with:
 tofu_version: 1.6.0

- name: OpenTofu Init
 run: tofu init
```

## **5. Team Communication**

### **Communication Checklist:**

- Notify team of migration completion
- Share migration results and findings
- Document lessons learned
- Update team procedures
- Schedule knowledge sharing session
- Provide migration support resources

```
OpenTofu Migration Completion Report

Project: [Project Name]
Date: [Date]
Performed By: [Name]

Summary
Successfully migrated [Project Name] from Terraform v1.5.7 to OpenTofu v1.6.0.

Results
- ✓ State migration completed without errors
- ✓ Zero infrastructure changes detected
- ✓ All modules functioning correctly
- ✓ CI/CD pipelines updated
- ✓ Documentation revised

Validation
- Plan output: No changes
- Resource count: X resources unchanged
- State serial: Incremented correctly
- Outputs: All verified

Issues Encountered
- [None | List any issues and resolutions]

Next Steps
- Team training scheduled for [date]
- Monitoring for one week
- Additional projects to migrate: [list]

Rollback Plan
- State backup location: [path/URL]
- Rollback tested: [Yes/No]
- Rollback time estimate: [X minutes]
```

## **Rollback Procedures**

### **When to Rollback:**

- State file corruption detected
- Significant infrastructure drift
- Provider compatibility issues
- Critical functionality broken
- Team decision to postpone

### **Rollback Checklist:**

- Stop any active operations
- Restore state file from backup
- Restore configuration from Git tag
- Reinitialize with Terraform
- Validate restoration with `terraform plan`
- Document rollback reason
- Plan remediation

## Rollback Commands:

---

```
bash
```

```
Stop current operations
```

```
Press Ctrl+C if anything running
```

```
Restore state from backup
```

```
cp terraform.tfstate.pre-migration terraform.tfstate
```

```
cp terraform.tfstate.backup.pre-migration terraform.tfstate.backup
```

```
Restore configuration from Git
```

```
git checkout pre-migration-$(date +%Y%m%d)
```

```
Remove OpenTofu artifacts
```

```
rm -rf .terraform .terraform.lock.hcl
```

```
Reinitialize with Terraform
```

```
terraform init
```

```
Verify restoration
```

```
terraform plan
```

```
Should show same state as before migration
```

## Summary

This migration planning tutorial covers the essential knowledge needed to successfully migrate from Terraform to OpenTofu:

1. **Understanding Why:** Licensing changes, open-source continuity, and business risk mitigation
2. **Compatibility Assurance:** HCL, providers, state files, and modules all work identically
3. **Dependency Mapping:** Identifying all components that need validation during migration
4. **Systematic Approach:** Comprehensive checklists for every phase of migration

## **Hands-on Migration Lab - VPC + EC2 Migration**

# Lab Overview

**Objective:** Migrate a production-like AWS infrastructure project from Terraform to OpenTofu, validating state consistency and resource integrity throughout the process.

## Infrastructure Components:

- Multi-AZ VPC with public and private subnets
- Internet Gateway and NAT Gateway
- Route tables and associations
- EC2 instance with SSH access
- Security groups with minimal access
- Elastic IP for NAT Gateway
- SSH key pair for instance access

## Learning Outcomes:

- Experience a complete migration workflow**
- Validate state consistency before and after migration**
- Troubleshoot common migration issues**
- Document migration results**
- Understand rollback procedures**